

De-vectorization

- ❑ Cache space and bandwidth are scarce resources
- ❑ Compilers know this – but sometimes they have to store data that does not need to be stored.
- ❑ This impacts:
 - ❑ bandwidth
 - ❑ cache capacity
 - ❑ instruction scheduling

De-vectorization

- ❑ A typical problem with scratch data
- ❑ Difficult/impossible for the compiler to detect
- ❑ Depends on coding style

De-vectorization – Example

```
COMMON /SCRATCH/TMP (N)
```

```
DO I = 1, N  
    TMP (I) = ...
```

```
:
```

```
    ... = TMP (I)  
END DO
```

```
...
```

```
DO I = 1, N  
    TMP (I) = ..  
    :
```

```
END DO
```

- ❑ *Because $TMP()$ is global, the compiler has to store it in the first loop*
- ❑ *In the second loop, $TMP()$ is overwritten, but the compiler will most likely not see this*
- ❑ *The programmer may know that $TMP()$ is a scratch array only*

De-vectorization – Solutions

Array TMP needed later on:

```

REAL  T1, TMP (N)

DO I = 1, N
    T1 = ...
    :
    ... = T1
END DO
    :
    .

DO I = 1, N
    TMP (I) = ..
    :
END DO
  
```

Array TMP not needed later on:

```

REAL  T1

DO I = 1, N
    T1 = ...
    :
    ... = T1
END DO
    :
    .

DO I = 1, N
    T1 = ..
    :
END DO
  
```

Stripmining

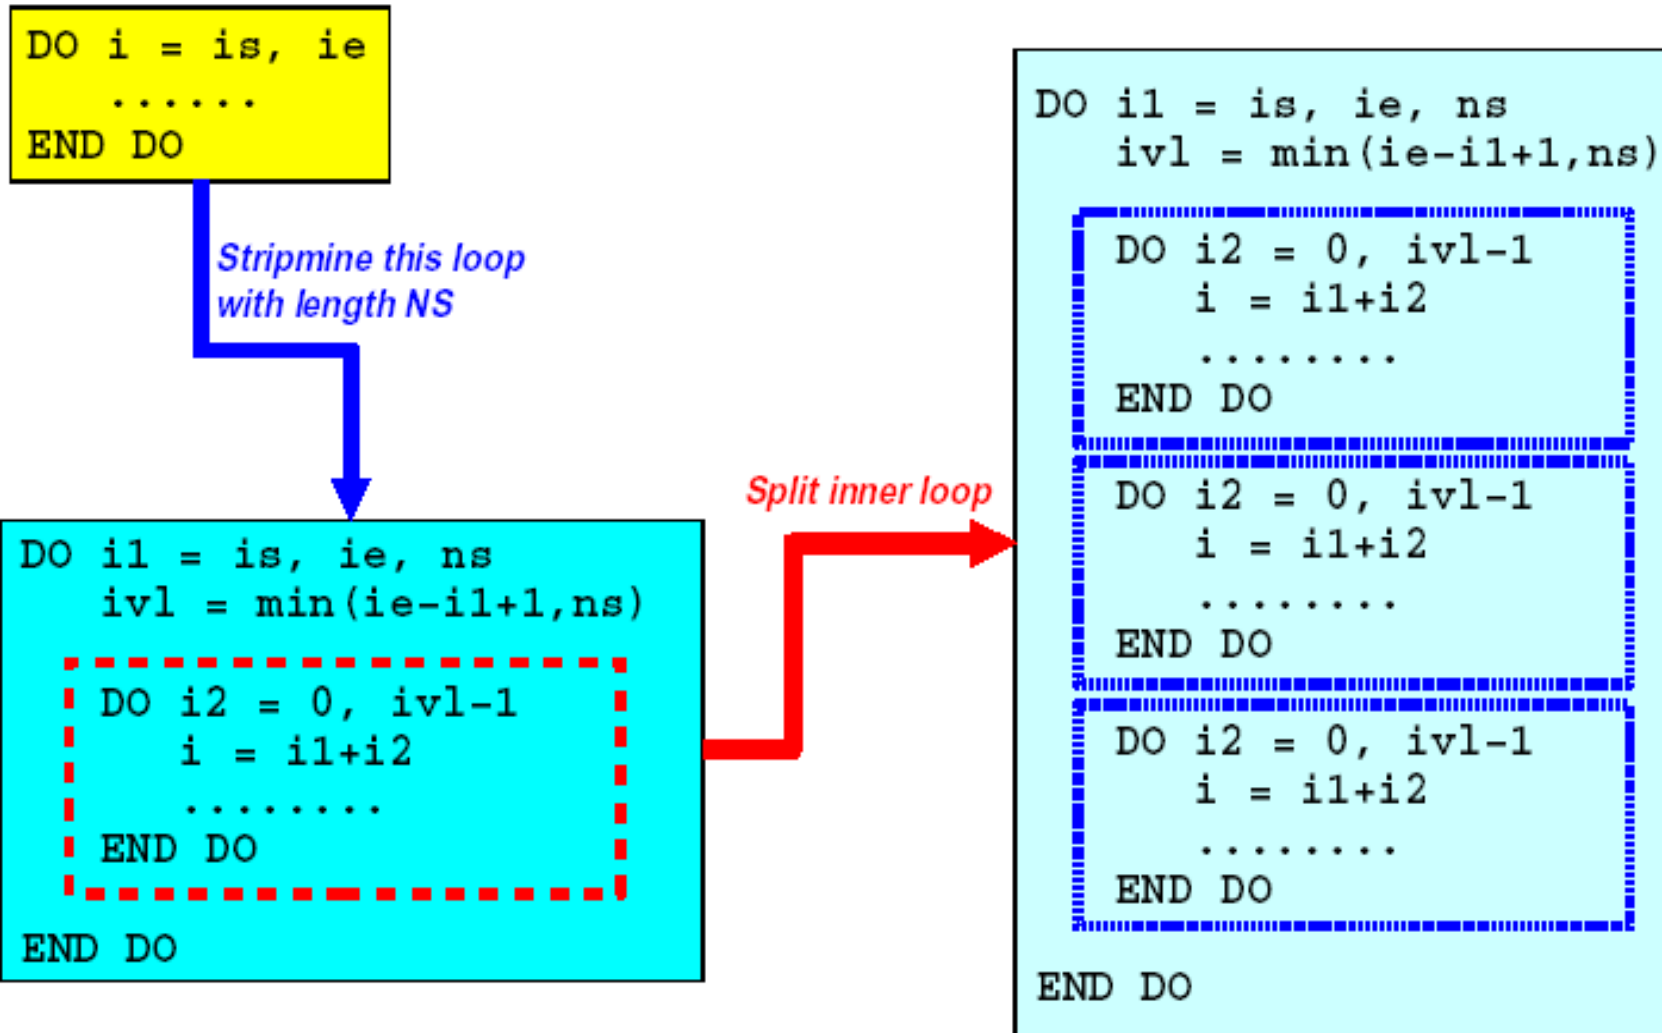
- ❑ *Large loops are difficult to optimize*
- ❑ *Especially the **register allocation** in the compiler has a hard time and can get confused*
- ❑ *Splitting the loop into smaller loops may improve performance*
- ❑ *However, this may cause scalars (local to the loop) to be replaced by vectors*
- ❑ *On very large loops this will increase memory usage notably*
- ❑ *Through **stripmining** memory usage can be kept under control*

```
DO I = 1, LONG  
  X(I) = ...  
  A = ...  
  Y(I) = A + ...  
END DO
```

Split loop in
two parts

```
DO I = 1, LONG  
  X(I) = ...  
  VA(I) = ...  
END DO  
  
DO I = 1, LONG  
  Y(I) = VA(I) + ...  
END DO
```

Stripmining – Code structure



Best practice

It is up to you to write code such that the compiler can find opportunities for optimization:

- ❑ Write efficient, but clear code
- ❑ Avoid very “fat” (bulky) loops
- ❑ Design your data structures carefully
- ❑ Minimize global data

Best practice

- ❑ Branches:
 - ❑ simplify where possible
 - ❑ try to split the branch part out of the loop
- ❑ Avoid function calls in loops (use inlining)
- ❑ Leave the low level details to the compiler

Summary

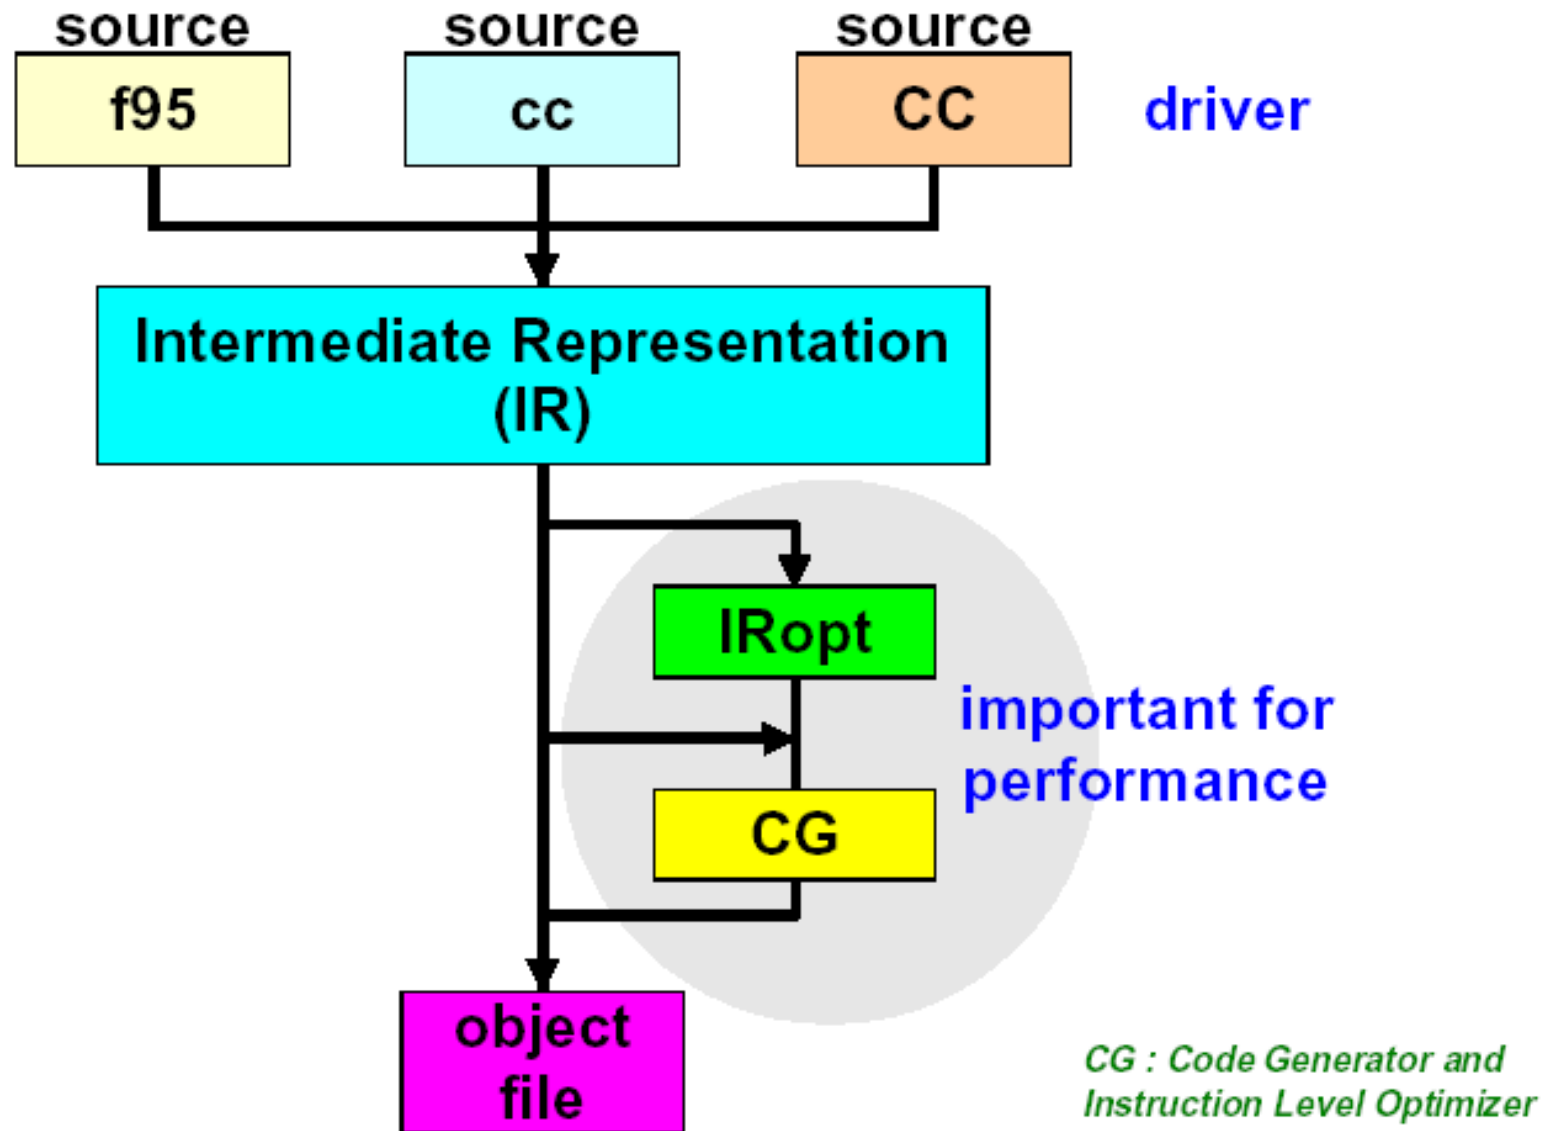
- ❑ Most tuning techniques presented here are generic, i.e. they (probably/hopefully) improve your code on all cache based systems.
- ❑ The *tuning parameters* may be different, though, since they depend on the underlying hardware:
 - ❑ cache sizes and levels
 - ❑ prefetch
 - and your problem's *memory footprint*
- ❑ Use the *best* compiler available on your platform.

The Sun Studio Compilers

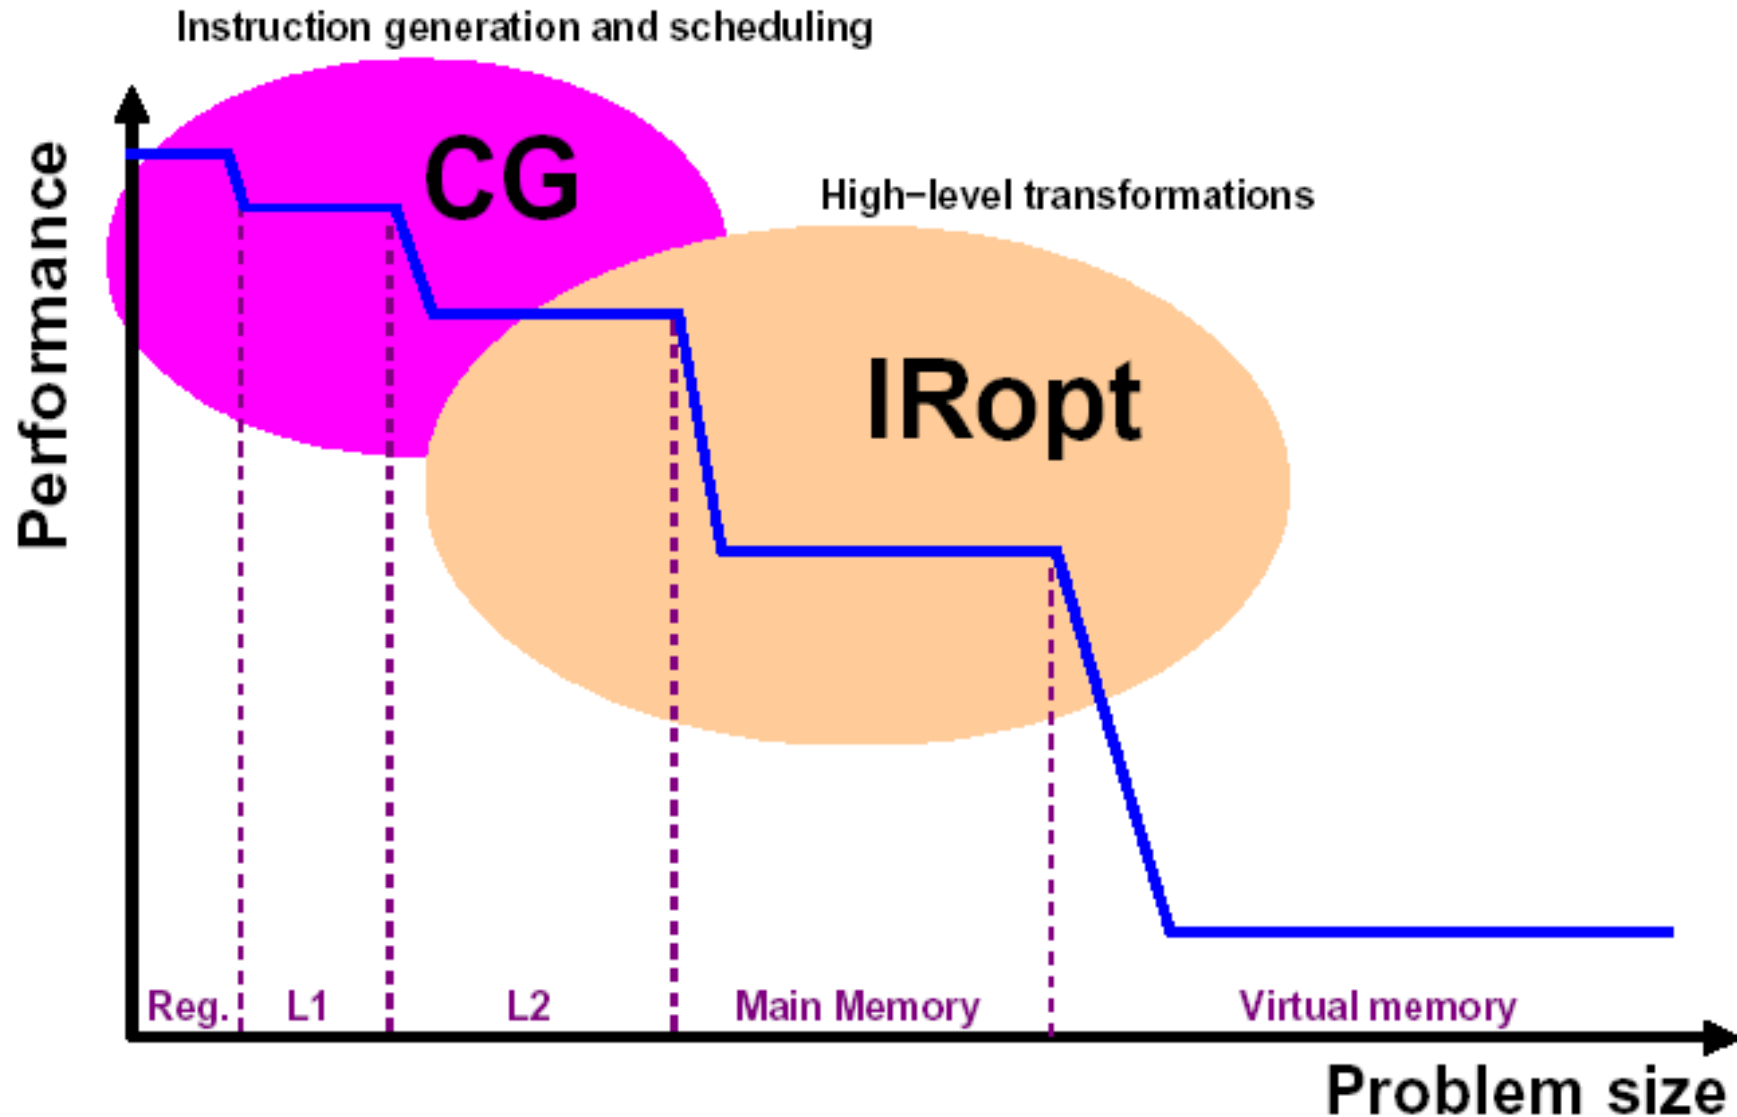
Sun Studio: Overview

- ❑ Compiler Components
- ❑ Compiler Options
- ❑ Compiler Commentary
- ❑ Best practice

Sun Studio: Compiler Components



Sun Studio: Who does what?



Sun Studio: Minimal Compiler Options

In general, one gets very good performance by just using 3 options for compiling and linking:

`-g -fast -m32`

(32-bit addressing)

`-g -fast -m64`

(64-bit addressing)

For the UltraSPARC-IV processor:

`-fast -xchip=ultra4 -m32`

(32-bit addressing)

`-fast -xchip=ultra4 -m64`

(64-bit addressing)

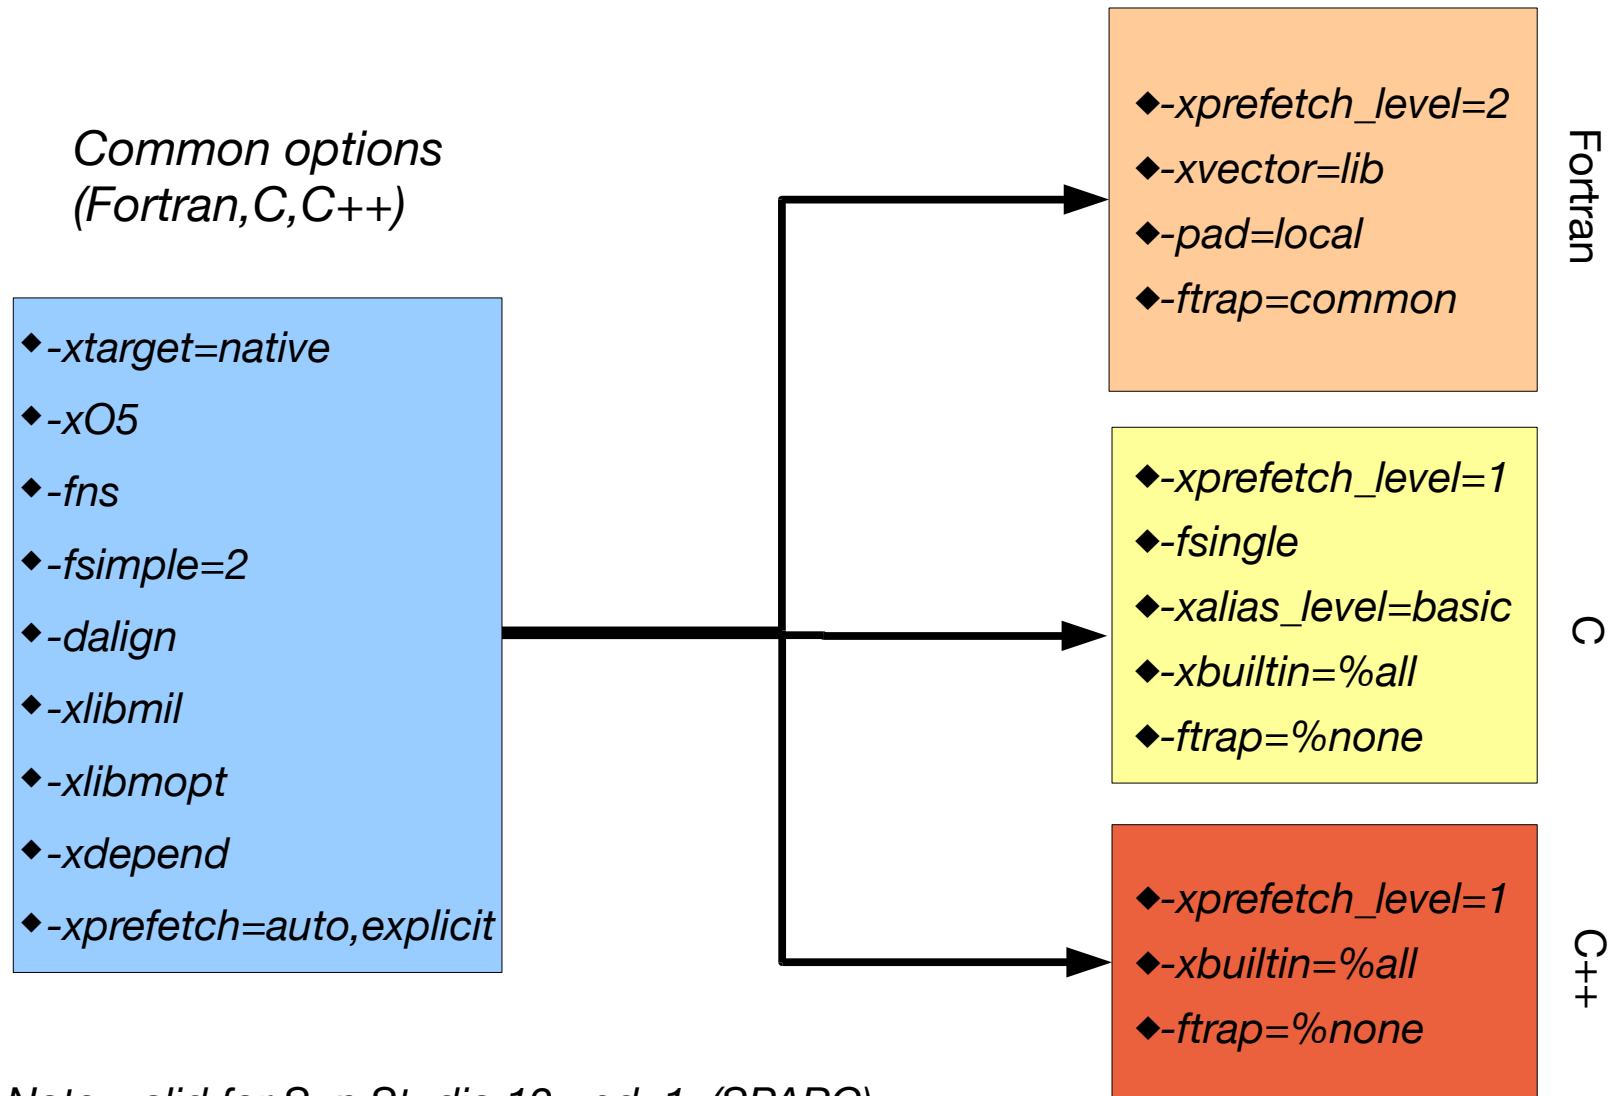
For other UltraSPARC processors:

US-T2: `-xchip=ultraT2`

US-IIIi: `-xchip=ultra3i`

US-IV+: `-xchip=ultra4plus`

Sun Studio: The -fast macro



Note: valid for Sun Studio 12 upd. 1 (SPARC)

Sun Studio: Recommendations

- ❑ Use a Makefile and `make` for compiling/linking:

```
OPT      = -g -fast
ISA      = -m32
CHIP     = -xchip=ultra4plus
CFLAGS   = $(OPT) $(ISA) $(CHIP)
```

- ❑ *Always start with `-fast` !*
- ❑ The compilers follow the 'rightmost option wins' rule, i.e. one can overrule options defined by the `-fast` macro.

Sun Studio: Recommendations

- ❑ `-fast` is a convenience macro that (in general) gives optimal performance with one single option
- ❑ `-fast` can change from one release to another!
- ❑ Use '`-fast -xdryrun`' to check what `-fast` expands to

What about other compilers? - I

- ❑ Other compilers have similar options, that combine many optimizations into a single option
- ❑ GCC: -O, -O3
- ❑ Intel: -O2 (default!), -O3, -fast
- ❑ look up in the manpages/documentation, what that corresponds to
- ❑ my experience: difficult to find out, not so easy to switch off 'unwanted options'

What about other compilers? - II


- ❑ Other compilers work in a similar way, but ...
- ❑ I've chosen Sun Studio for various reasons:
 - ❑ my own experience
 - ❑ it is easy to use
 - ❑ it makes it easy to test effects of certain options
 - ❑ it is a nice tool for teaching
 - ❑ it makes it easier for you to understand certain aspects in this class

Sun Studio: Optimization levels

- ❑ The level of optimization can be specified with the `-xOn` option
- ❑ Each level includes the lower levels
- ❑ Choices for `n`:
 - ❑ `n = 1` – Basic block level optimizations
 - ❑ `n = 2` – Some additional optimizations
 - ❑ `n = 3` – Loop transformations and modulo scheduling
 - ❑ `n = 4` – Intra-file inlining and pointer tracking
 - ❑ `n = 5` – Aggressive optimizations
- ❑ Use the `-xipo` option to get intra- **and** interfile inlining (requires minimum `-xO4`)

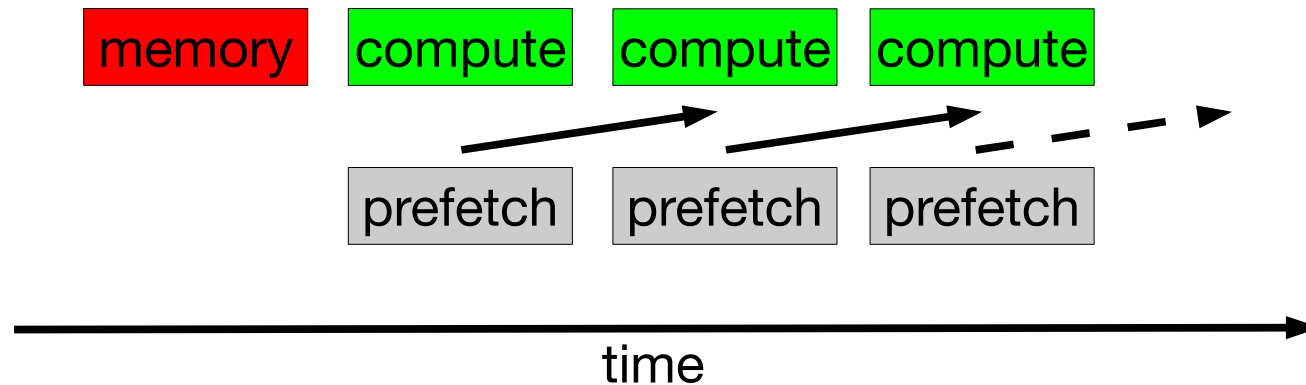
Inter Procedural Optimization

- ❑ With the -xipo option, the compiler stores additional information into the object files
- ❑ This information is used during the link phase to perform additional optimizations



Prefetch: Hiding memory latency

- ❑ The number of clock cycles to access memory increases with the CPU clock speed.
- ❑ Prefetch is a way to overcome this:
 - Fetch data ahead in time, anticipating future use.
- ❑ Special prefetch instructions must be available



Prefetch Support


- ❑ Prefetch is a common feature in modern CPUs: both data and instruction prefetch.
- ❑ Implementation is system dependent!
- ❑ Examples of CPUs that support prefetch:
 - ❑ UltraSPARC IIIi, IIICu, IV, IV+
 - ❑ AMD Opteron
 - ❑ Intel (since Pentium 3, 4)
- ❑ Compilers have to support prefetch.
- ❑ There is hardware prefetch, too

Sun Studio: Prefetch options

- ❑ Automatic – done by the compiler:
 - ❑ -xprefetch=yes
 - ❑ Control level and type:
 - ❑ -xprefetch_level=n (n = 1, 2 or 3)
 - ❑ -xprefetch_auto_type=[no%]indirect_array_access
- ❑ Explicit – controlled by the user:
 - ❑ Either through functions calls (C/C++) or directives (Fortran).
 - ❑ Can be combined with automatic prefetch.
 - ❑ For more information see the compiler manual.

Prefetch: example

Example: Matrix times vector in C (row version)



US-IIIi @ 1062 MHz
L1 : 64 kB
L2 : 1 MB
Peak : 2.1 Gflop/s

Pointer overlap

```
void vecadd(int n, double *a, double *b, double *c)
{
    for(int i = 0; i < n; i++)
        c[i] = a[i] + b[i];
}
```

```
vecadd(n, &a[0], &b[0], &a[1]);
```

```
void vecadd(n, &a[0], &b[0], &a[1])
{
    for(int i = 0; i < n; i++)
        a[i+1] = a[i] + b[i];
}
```

**Data
dependency!**

Sun Studio: C code and `-xrestrict`


- ❑ Pointer aliasing problem: The C compiler has to assume that different pointers may overlap:
 - ❑ Correct – but non-optimal – code will be generated
 - ❑ Only the programmer might know, that there is no overlap.
- ❑ You can tell the compiler that there is no overlap, with the `-xrestrict` option.
- ❑ Note: It is then your responsibility that this assumption will not be violated!

Sun Studio: C code and -xrestrict

- ❑ Syntax: -xrestrict=n, where n is one of
 - ❑ %none : all pointers may overlap (default)
 - ❑ %all : no pointers overlap
 - ❑ f1[,f2] : a comma separated list of function names
- ❑ -xrestrict is the same as -xrestrict=%all
- ❑ Warning: Use -xrestrict only if you can make sure that the pointers never overlap!

Sun Studio: C code and -xrestrict

Vector addition example: non-overlapping vectors!



US-III @ 900 MHz
L1 : 64 kB
L2 : 8 MB
Peak : 1.8 Gflop/s
Sun Studio 10

Pointer overlap – fix your code

- ❑ If you can assure that the pointers don't overlap, you can fix your code using the C99 'restrict' keyword:

```
void  
vecadd(int n, double * restrict a,  
        double * restrict b, double * restrict c)  
{  
    for(int i = 0; i < n; i++)  
        c[i] = a[i] + b[i];  
}
```

- ❑ Advantage: local change, no global assumption
- ❑ Needs a C99 compiler to be portable!

Fortran: The -stackvar option

```

program stackvar

do i=1, 5
  call mysub()
enddo

stop
end

subroutine mysub()
k=k+1
print *, 'k = ', k
return
end

```

What is wrong here?

```

% f95 -x03 stack.f95
% ./a.out
k = 1
k = 2
k = 3
k = 4
k = 5

```

Intended(?) behaviour

```

% f95 -x03 -stackvar stack.f95
% ./a.out
k = 1
k = 1
k = 1
k = 1
k = 1

```

With -stackvar –
but still wrong!

The code is not initializing k in mysub(), i.e. the behaviour is compiler dependend!
Use -xcheck=init_local to find such problems.

Sun Studio: Useful options – 1

- ❑ -flags : Lists all the available compiler flags on the screen (long list)
- ❑ -xhelp=readme : Displays the README file (release notes) on the screen
- ❑ -xdryrun : see what the compiler would do (macro expansion, no compilation!)
- ❑ -V : Shows the compiler version

```
% f90 -V
f90: Sun Fortran 95 8.4 SunOS_sparc 2009/06/03
Usage: f90 [ options ] files.  Use 'f90 -flags' for details
```


Sun Studio: Useful options – 2

- ❑ -g : Generates debugging information and adds compiler commentary to the object file.
Necessary if the performance analysis tools should be used.
- ❑ **Note:** -g can be used in connection with optimization!
- ❑ -xlibmopt : Optimized version of libm
- ❑ -xlibmil: Use fast assembly language inline templates for standard functions
- ❑ -xvector : vectorized intrinsics, e.g. sqrt()

Sun Studio: Useful options – 3

- ❑ `-xprefetch_level=n` : Try different levels of prefetching ($n = 1, 2, 3$)
- ❑ `-xrestrict` : (C/C++) tell the compiler that your pointers do point to 'restricted areas' in memory, i.e. no overlap of arrays
- ❑ `-stackvar` : (Fortran) keep local data on stack
- ❑ `-xlic_lib=sunperf` : Link with the Sun Performance Library (LAPACK, BLAS, Sparse-BLAS (Netlib and NIST), FFT, Sparse Solvers)

Sun Studio: Useful options – 4

- ❑ Options that are useful under development – but should be avoided in production:
 - ❑ -C : array boundary checking (Fortran)
 - ❑ -pg/-p : Unix profilers – you don't really need them, use the Performance Analyzer instead.
 - ❑ -xcheck=... : implements runtime checks.
 - ❑ -xcheck=init_local : initialize local variables with a value that is likely to cause an arithmetic exception (Fortran)
 - ❑ -xcheck=stkovf : check for stack overflow.
 - ❑ -Xlist : check your Fortran code for inconsistencies (commented source in .lst file)

Sun Studio: On Intel/AMD as well

- ❑ For Solaris 10 on x86 and x86_64 (x64)
 - ❑ Sun Studio 12 update 1 IDE, compilers & tools
 - ❑ Sun Studio Express compilers & tools (Technology Preview) – if available
- ❑ For Linux (i386 and AMD64):
 - ❑ Sun Studio 12 update 1 IDE, compilers & tools
 - ❑ Sun Studio Express compilers & tools (Technology Preview) – if available

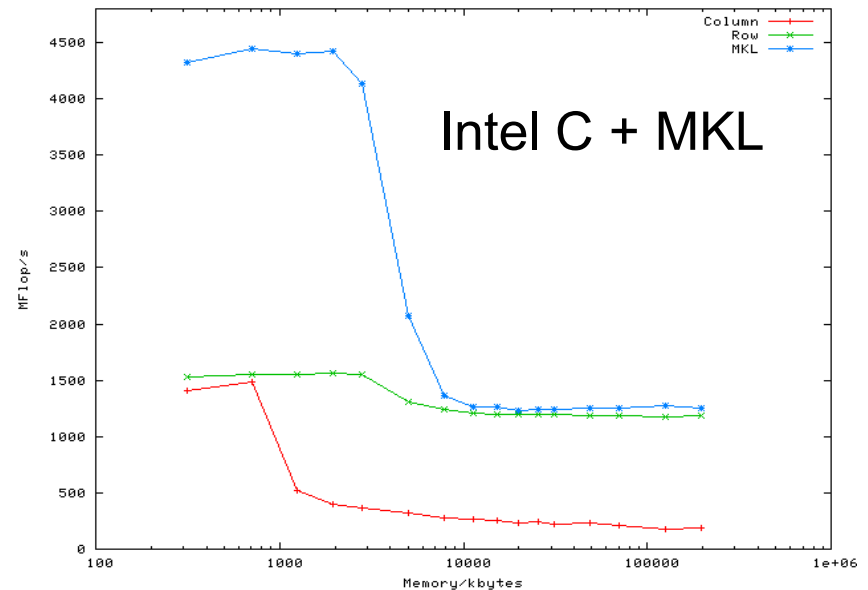
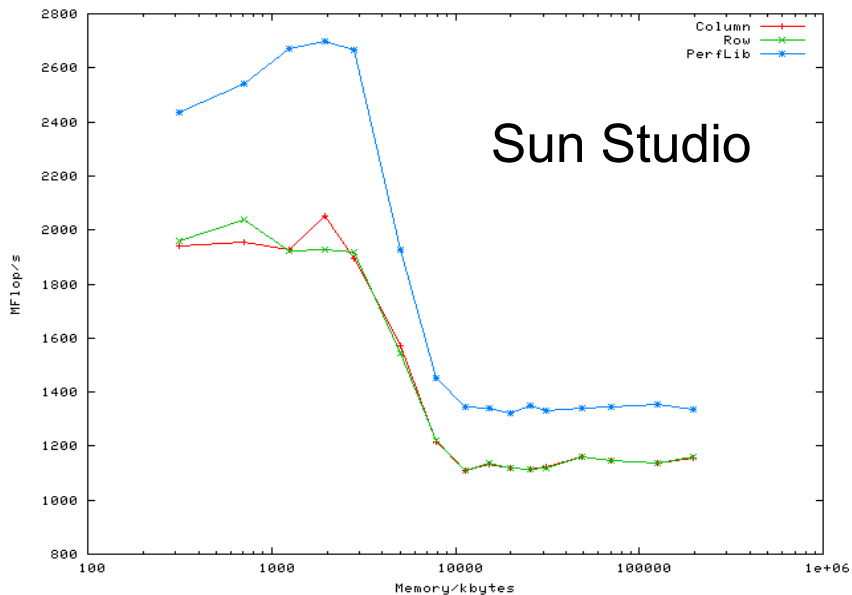
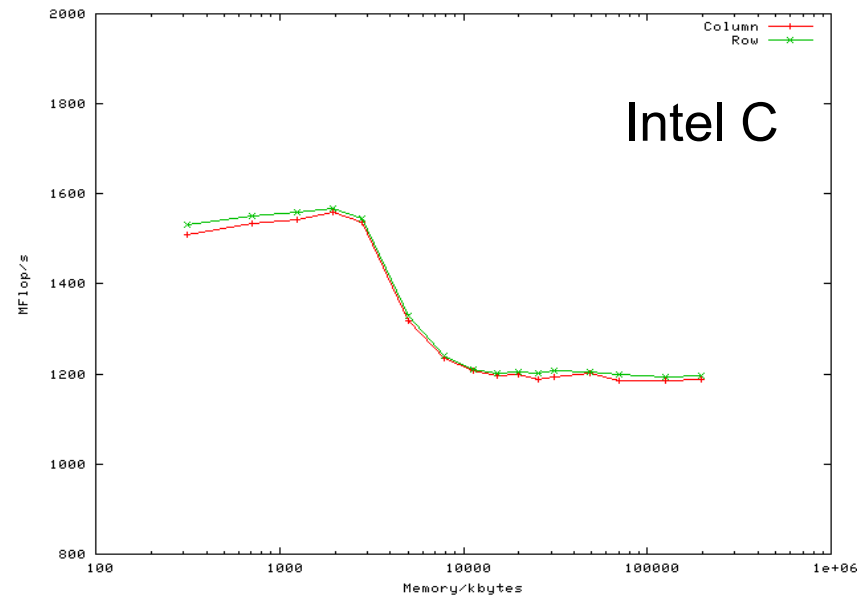
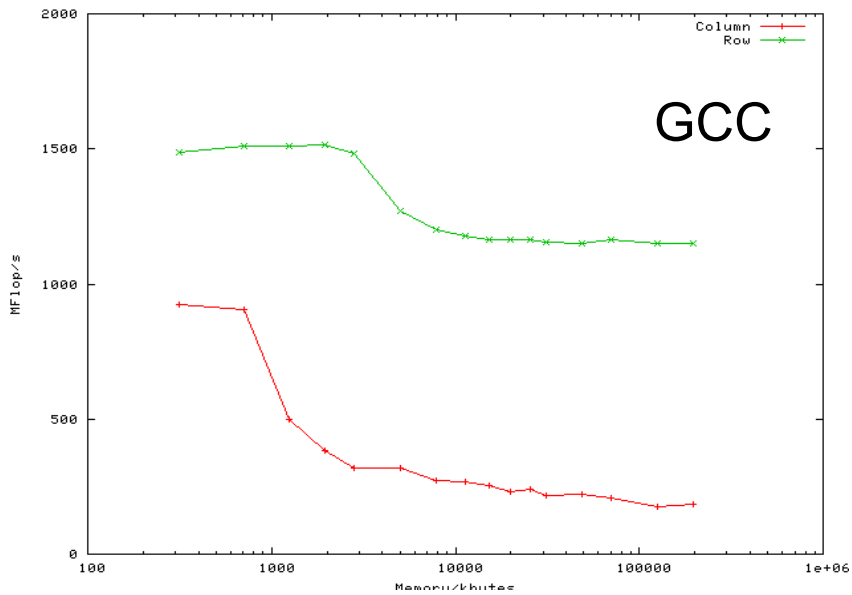
Sun Studio: Where to get it?

- ❑ Go to <http://developers.sun.com/sunstudio>
- ❑ Download and install it – it is for free!
- ❑ Note: Check if there are patches available, and install them as well
 - ❑ <http://developers.sun.com/sunstudio/downloads/patches/>
- ❑ Best practice: It is always good to have more than one compiler at hand!
 - ❑ to test compatibility / portability
 - ❑ to exclude compiler bugs

Matrix times vector


- ❑ 2 code versions: **Column** and **Row**
- ❑ 3 compilers under Linux:
 - ❑ GCC
 - ❑ Intel C
 - ❑ Sun Studio C
- ❑ 2 libraries:
 - ❑ Intel MKL
 - ❑ Sun Performance Library
- ❑ Note: this is a quick comparison! Based on versions available in June 2008!

Matrix times vector



Matrix times vector

Comparison for large data sets:



Ease of use – vendor library

Intel/GCC with Intel MKL:

- ❑ several compiler options
- ❑ several linker options
- ❑ depends on platform: IA32 or EMT64
- ❑ set up of run-time environment
- ❑ different downloads, installations and licenses

❑ Sun Studio:

- ❑ one linker option: `-xlic_lib=sunperf`
- ❑ no need to set up run-time environment
- ❑ one download, installation – and no license

Compiler commentary

Sun Studio: Compiler Commentary

How do I know what the compiler did with my code?

- ❑ Compile with -g (C++: -g0) and use the `er_src` command on the object files:
 - ❑ This generates a code listing with comments from the compiler (also used in the Performance Analyzer GUI).
- ❑ Command: `er_src -src foo.c foo.o`
- ❑ Note: the examples on the next slides are from the Solaris SPARC version of Sun Studio – the x64 versions can differ

Sun Studio: Compiler Commentary

```
// ex.c
#define A(i,j) a[(i)*n + (j)]
#define B(i,j) b[(i)*n + (j)]
#define C(i,j) c[(i)*n + (j)]

void
mysub(int m, int n, double *a, double *b,
      double *c) {

    int i, j;

    for(j = 0; j < n; j++)
        for(i = 0; i < m; i++)
            C(i,j) = A(i,j) + B(i,j);
}
```

Sun Studio: Compiler Commentary

```
5. void
6. mysub(int m, int n, double *a,
        double *b, double *c){
```

<Function: mysub>

```
7.
8.     int i, j;
9.
```

Source loop below has tag L1

```
10.     for(j = 0; j < n; j++)
```

Source loop below has tag L2

L2 scheduled with steady-state cycle count = 5

L2 unrolled 2 times

L2 has 2 loads, 1 stores, 6 prefetches, 1 FPaddds, 0
FPMuls, and 0 FPdivs per iteration

```
11.         for(i = 0; i < m; i++)
12.             C(i,j) = A(i,j) + B(i,j);
13.     }
```

```
% cc -g -fast -c ex.c
% er_src -src ex.c ex.o
```

Sun Studio: Compiler Commentary

```
% cc -g -fast -xrestrict -c ex.c
% er_src -src ex.c ex.o
```

```
5. void
6. mysub(int m, int n, double *a,
        double *b, double *c){
```

```
<Function: mysub>
```

```
7.
8.     int i, j;
9.
```

Source loop below has tag L1

L1 interchanged with L2

L1 scheduled with steady-state cycle count = 2

L1 unrolled 4 times

L1 has 2 loads, 1 stores, 3 prefetches, 1 FPaddds, 0
FPMuls, and 0 FPDivs per iteration

```
10.     for(j = 0; j < n; j++)
```

Source loop below has tag L2

L2 interchanged with L1

```
11.         for(i = 0; i < m; i++)
12.             C(i,j) = A(i,j) + B(i,j);
13.     }
```

Sun Studio: Compiler Commentary

```
% cc -g -fast -xrestrict -xunroll=1 -c ex.c  
% er_src -src ex.c ex.o
```

```
5. void  
6. mysub(int m, int n, double *a,  
        double *b, double *c){  
    <Function: mysub>  
7.  
8.     int i, j;  
9.
```

Source loop below has tag L1

L1 interchanged with L2

```
10.     for(j = 0; j < n; j++)
```

Source loop below has tag L2

L2 interchanged with L1

```
11.         for(i = 0; i < m; i++)  
12.             C(i,j) = A(i,j) + B(i,j);  
13.     }
```

Sun Studio: Compiler Commentary

```
% cc -g -fast -xrestrict -c ex2.c
% er_src -src ex2.c ex2.o
```

```
7. void
8. mysub(int m, int n, double *a,
        double *b, double *c){
```

<Function: mysub>

```
9.
10.     int i, j;
11.
```

Source loop below has tag L1

```
12.     for(j = 0; j < n; j++)
```

Source loop below has tag L2

L2 could not be pipelined because it contains calls

```
13.         for(i = 0; i < m; i++) {
14.             C(i,j) = A(i,j) + B(i,j);
```

Function mean not inlined because the compiler has not seen the body of the routine

```
15.             mean(C(i,j));
16.         }
17. }
```


Compiler Commentary – I

Why is it useful?

- ❑ Compilers are no longer a black box!
- ❑ What the compiler has done – or hasn't done/couldn't do – to the code is made visible to the programmer.
- ❑ Useful information is provided, so a programmer can take action, e.g.
 - ❑ code changes
 - ❑ different set of compiler options

Compiler Commentary – II

More reasons, why it is really useful:

- ❑ annotations are where they belong: in the code – and not on the screen during compilation (or in a log file)
- ❑ review at any time – even a long time after the code has been compiled
- ❑ visible in the analyzer output as well, together with runtime profile

Other tricks

Reconstruct the compiler options from the object files and/or executable:

- ❑ dwarfdump file.o (C/C++/Fortran)
- ❑ dumpstabs file.o (Fortran, before Studio 12u1)
- ❑ and look for
 - ❑ command_line (dwarfdump)
 - ❑ CMDLINE (dumpstabs) orin the output.
- ❑ Very useful to check what has been done.
- ❑ Note: dwarfdump works for GCC .o files, too

GCC vs Sun Studio:

Get the best of both worlds

- ❑ Problem (typical with C++): my code doesn't compile with Sun Studio, due to “GCC extensions” - but I am bound to Solaris SPARC
- ❑ Solution: Get GCCFSS from <http://cooltools.net>
- ❑ What is it?
 - ❑ A GCC front-end (understands all the GCC'isms)
 - ❑ plus the optimizing Sun Studio backend (gives you full performance)
 - ❑ understands Studio options: `gcc -O3 -fast ...`

Analysis tools

Analysis tools

- ❑ analysis tools are useful to detect bottlenecks in codes
- ❑ modern analysis tools (unlike “old” profilers) work even on 'non-instrumented' code: no need to recompile (in principle)
- ❑ runtime profiles down to the source level (profilers usually work on function/subroutine level)

Analysis tools

- ❑ Sun: Sun Studio Performance Analyzer
 - ❑ Solaris (SPARC and Intel/AMD)
 - ❑ Linux
- ❑ Intel: Vtune Performance Analyzer (Windows/Linux)
- ❑ AMD: CodeAnalyst (Windows/Linux)
- ❑ Mac OS X: Shark (part of Xcode)
- ❑ more ... ?

Sun Studio: Performance Analyzer

- ❑ Sun Studio provides a powerful toolset for runtime analysis
- ❑ Both GUI and command line tools
 - ❑ analyzer – GUI for collecting and analyzing performance data
 - ❑ collect – Command to collect performance data
 - ❑ er_print – Command to analyze performance data in ASCII format (good for scripting)

Sun Studio: Performance Analyzer

The screenshot shows the Sun Studio Performance Analyzer window for a test named 'test.1.er'. The 'Functions' tab is active, displaying a table of functions with their User CPU times. The 'matmult' function is selected, and its details are shown in the right-hand pane.

User CPU (sec.)	User CPU (sec.)	Name
3.178	3.178	<Total>
1.716	1.716	cblas_dgemm
1.408	1.408	matmult
0.022	0.022	erand48_r
0.022	0.055	init_matrix
0.011	0.033	drand48
0.	1.716	main

Data for Selected Object:

- Name: matmult
- PC Address: 2:0x000008E0
- Size: 753
- Source File: ./multiply.c
- Object File: ./matmult
- Load Object: <matmult>
- Mangled Name:
- Aliases:

Process Times (sec) / Counts

	Exclusive	Inclusive
User CPU:	1.408 (44.29%)	1.408 (44.29%)

Sun Studio: Performance Analyzer

The screenshot displays the Sun Studio Performance Analyzer window for a test file named 'test.1.er'. The interface includes a menu bar (File, View, Timeline, Help), a toolbar, and a 'Find Text' search field. The main window is divided into several panes:

- Functions Pane:** Lists 'User CPU' and 'CPU (sec.)' for the selected function.
- Source File:** Shows the source code for 'multiply.c' located at '/home/bd/IMM/02613_E05/assign1/multiply.c'. The code includes standard headers, defines macros for array access, and contains a 'matmult' function. The function iterates over a 2D array 'C' using nested loops for 'u' and 'v', calculating the product of elements from arrays 'A' and 'B'.
- Timeline Pane:** Displays a timeline of events, with the 'matmult' function highlighted in green, indicating its execution time of 1.298 seconds.
- Summary Pane:** Provides detailed information about the selected object, including its name, PC address, size, source file, object file, load object, and aliases.
- Process Times (sec) / Counts:** A table showing the execution time and counts for the 'User CPU'.

Process Times (sec) / Counts			
		Exclusive	Inclusive
User CPU:	1.298 (40.83%)	1.298 (40.83%)	

Hardware Performance Counters

- ❑ Almost all modern CPUs have built-in hardware performance counters:
 - ❑ How many instructions were executed?
 - ❑ How many clock cycles were used?
 - ❑ How many L1 data cache misses occurred?
 - ❑ ...
- ❑ The supported counters are usually listed in the architecture reference manuals.
- ❑ Be aware: The counter names are not for beginners!

Using the Performance Counters

- ❑ Native Solaris tools:

- ❑ cputrack – track CPU usage on a per-process basis

- ❑ example:

- ```
% cputrack -c EC_ref,EC_misses -p <PID>
```

- ❑ cpustat, busstat – those require root privileges

- ❑ Sun Performance Analyzer tools:

- ❑ data sampled into a database

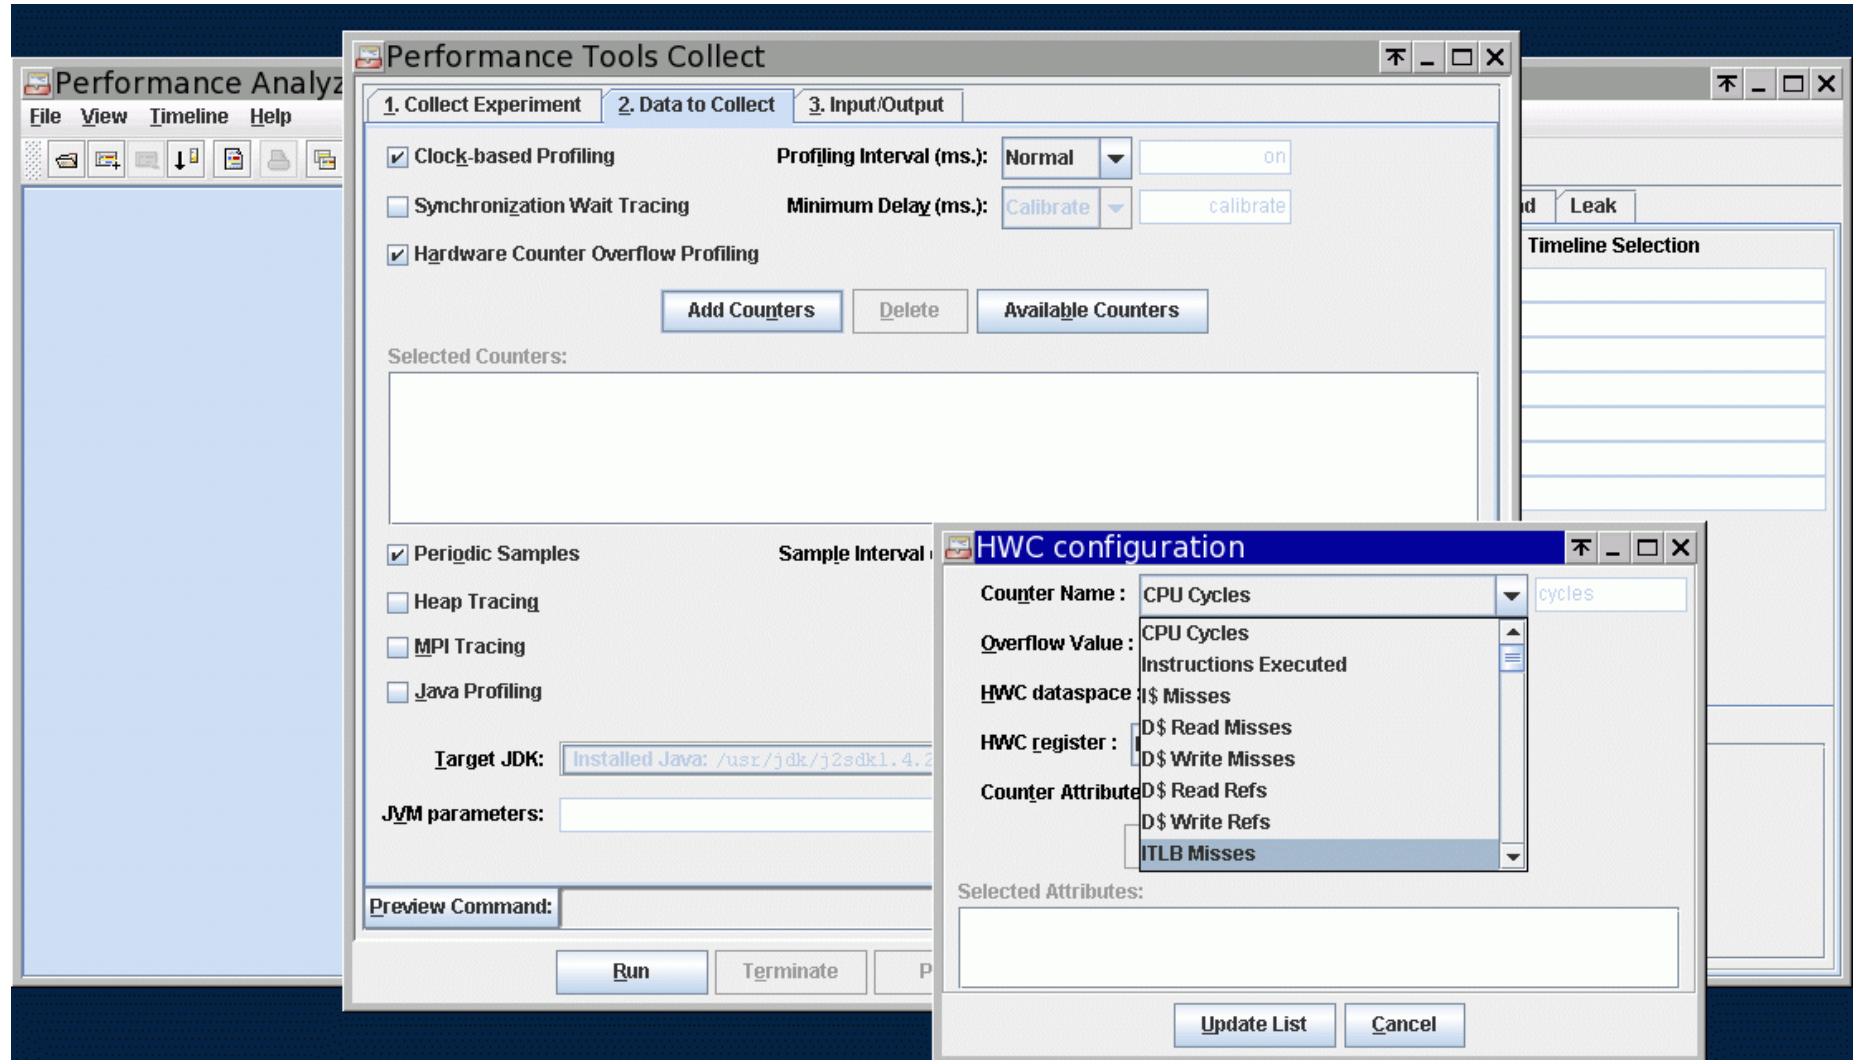
- ❑ information on source line level

# Using the Performance Counters

- ❑ Available performance counters:
  - ❑ system and CPU dependent
  - ❑ get a list:
    - ❑ `% cputrack -h`
    - ❑ `% collect` (no argument)
  - ❑ example: no. of available performance counters on
    - ❑ US-IIIi: 64
    - ❑ US-IV: 70
    - ❑ US-IV+: 101

# Using the Performance Counters

Activating performance counters in analyzer:





# Sun Studio: Performance Analyzer

Analyzer demo

# Tuning Guide – compact version

- ❑ Make a 'baseline' version (with different data sets/memory requirements)
- ❑ Try to find the best compiler options
  - ❑ with or w/o prefetching
  - ❑ ...
- ❑ Use analysis tools to locate the 'hot spots'
- ❑ Introduce code changes
- ❑ Repeat the last two steps until you are satisfied



End of lecture 2