

Application Tuning

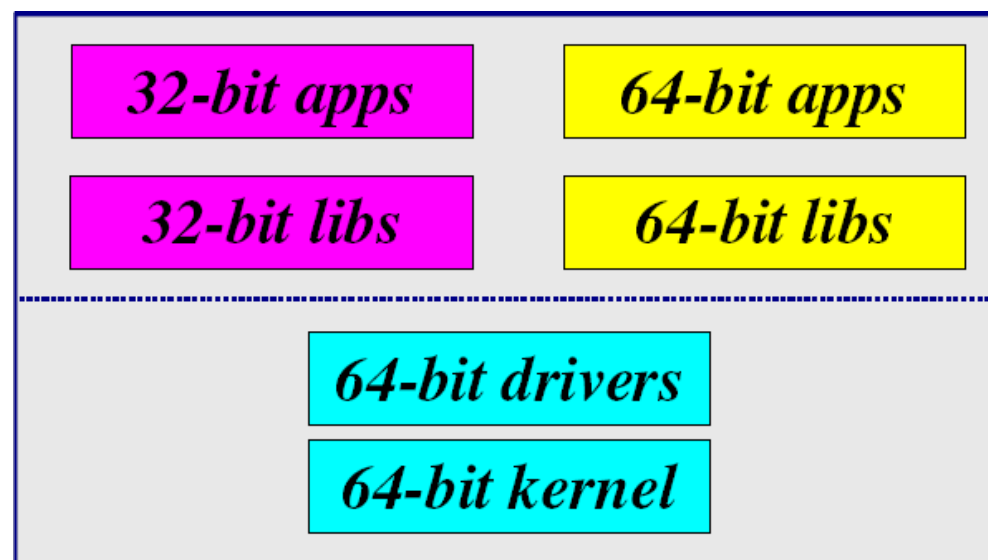
Selected Topics

Application Tuning


- ❑ Selected Topics:
 - ❑ 32- vs 64-bit
 - ❑ binary data portability
 - ❑ floating point numbers and IEEE 754
 - ❑ compiler options
 - ❑ case studies
 - ❑ large pages
- ❑ Summary

32-bit vs 64-bit issues

- ❑ 64-bit operating systems
- ❑ Implication: The address space of a single application can be larger than 4 GB



32-bit vs 64-bit issues



- ❑ **Addresses \neq Data**
- ❑ An 'n'-byte data type fills always n bytes in memory (byte addressable)
- ❑ I.e. the next element is n bytes further in memory
- ❑ This increment is not related to the size of the addresses (32-bit or 64-bit)

32-bit vs 64-bit issues

<u><i>C data type</i></u>	<u><i>ILP32</i></u> <i>(bits)</i>	<u><i>LP64</i></u> <i>(bits)</i>
<i>char</i>	8	<i>same</i>
<i>short</i>	16	<i>same</i>
<i>int</i>	32	<i>same</i>
<i>long</i>	32	64
<i>long long</i>	64	<i>same</i>
<i>pointer</i>	32	64
<i>enum</i>	32	<i>same</i>
<i>float</i>	32	<i>same</i>
<i>double</i>	64	<i>same</i>
<i>long double</i>	128	<i>same</i>

UNIX and Linux support LP64; Windows 64-bit uses LLP64, where long stays 32 bits

(p)ldd and LD_LIBRARY_PATH

How to check which shared-libraries are loaded?

- ❑ Static check: use the ldd command
 - ❑ `$ ldd executable`
- ❑ Dynamic check: use pldd on the PID
 - ❑ `$ pldd pid`
 - ❑ Solaris only
 - ❑ there are scripts available for Linux as well

(p)ldd and LD_LIBRARY_PATH

How to change the search path for dynamic libraries?

- ❑ Use LD_LIBRARY_PATH – but use it with care!
- ❑ Solaris can distinguish between 32- and 64-bit:
 - ❑ LD_LIBRARY_PATH – common
 - ❑ LD_LIBRARY_PATH_32 – for 32-bit apps
 - ❑ LD_LIBRARY_PATH_64 – for 64-bit apps
- ❑ Linux: only one setting !!!

(p)ldd and LD_LIBRARY_PATH

Best practice:

- ❑ Compile the path into your application:
 - ❑ Sun Studio: `-R <path_to_lib>`
 - ❑ GCC: `-Wl,-rpath <path_to_lib>`
 - ❑ `ld.so.1` will then use this path
- ❑ Avoid `LD_LIBRARY_PATH` in your shell environment – use a wrapper script for the application

Binary data storage

- ❑ Storing your data in binary format
- ❑ Advantages:
 - ❑ compact
 - ❑ fast
 - ❑ no loss of precision
- ❑ Drawbacks:
 - ❑ not “human readable”
 - ❑ data analysis more complicated
 - ❑ and ...

Binary data storage

- ❑ Example: integer 0x12345678 (hexadecimal)

```
value = 0x12345678;           // 305419896
printf("%d\n", value);
fwrite(&value, sizeof(value), 1, fptr);
```

- ❑ Write it ...

- ❑ ... on i386:

```
305419896
Architecture: i386
Value written to endian_i386.dat.
```

- ❑ ... on SPARC:

```
305419896
Architecture: sparc
Value written to endian_sparc.dat.
```

Binary data storage

❑ Read it:

```
fread(&value, sizeof(value), 1, fptr);  
printf("%d\n", value);
```

❑ on i386 data from i386:

Architecture: i386

Read from endian_i386.dat: 305419896

❑ on i386 data from SPARC:

Architecture: i386

Read from endian_sparc.dat: 2018915346




Little Endian vs Big Endian

- ❑ The order in which the bits are interpreted has not been standardized!
- ❑ Two 'popular' formats in use
 - ❑ Big Endian – SPARC, PowerPC, ...
 - ❑ Little Endian – Intel x86, AMD64, ...
- ❑ This is an issue when using the same binary data file on both platforms ...

Little Endian vs Big Endian

- Example: integer 0x12345678 (hexadecimal)



- Check with 'od' command:

```
$ od -x endian_sparc.dat
00000000 1234 5678
00000004

$ od -x endian_i386.dat
00000000 7856 3412
00000004
```

Little Endian vs Big Endian

- ❑ This is something you should be aware of when working with binary data!
- ❑ Tools:
 - ❑ Sun Fortran: -xfilebyteorder option
 - ❑ Portland Fortran compiler
 - ❑ swab() subroutine (low level)
 - ❑ ...

Floating point numbers & IEEE 754

- ❑ Remember: `-fast` expands to a set of options, and two of them are:
 - ❑ `-fns=yes`: faster – but non-standard – handling of floating-point arithmetic exceptions and gradual underflow (small numbers)
 - ❑ `-fsimple=2`: aggressive floating-point optimizations
- ❑ If your code requires to follow strictly the IEEE Standard for Binary Floating Point Arithmetic (IEEE 754), you can use:
 - ❑ `-fast -fns=no -fsimple=0` (or `-fsimple=1`)

Floating point numbers & IEEE 754

Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=0:

```
1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 17
L-1 has 1 loads, 1 stores, 2 prefetches, 0 FPadds,
0 FPMuls, and 1 FPdivs per iteration

Floating point numbers & IEEE 754

Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=2:

```
1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.
Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 1
L-1 unrolled 8 times
L-1 has 1 loads, 1 stores, 4 prefetches, 0 FPadds,
      1 FPMuls, and 0 FPdivs per iteration
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```


Floating point numbers & IEEE 754

Effects of `-fns=[yes|no]` - a case study:

- ❑ Multiplying a Toeplitz matrix with a random matrix in Matlab took about 20 times longer than the product of two equally sized random matrices (on SPARC).
- ❑ This didn't happen on the Linux/Intel platform – here both operations took approximately the same time.
- ❑ An investigation of the structure of the Toeplitz matrix showed, that it had a large number of entries with *subnormal* numbers, i.e. numbers smaller than $10E-300$.

Floating point numbers & IEEE 754

One slice of the Toeplitz matrix:



Floating point numbers & IEEE 754

- ❑ What is going wrong here?
- ❑ Explanation:
 - ❑ The operations with the subnormal numbers result in lots of gradual underflows, every one causing a hardware trap on the SPARC platform. Those traps are really expensive (pipeline flushes, etc).
- ❑ Why's that?
 - ❑ Matlab on SPARC is compiled without the optimization option -fns that flushes those small numbers to zero.

Floating point numbers & IEEE 754

- ❑ Runtime of the Matlab native version with a 500x500 Toeplitz matrix: 14.45 secs
- ❑ Used the Matlab compiler mcc (mcc calls cc from Sun Studio) with the right optimization option (-fns=yes) in the mbuildopt.sh file.
- ❑ The runtime of the same example was reduced to 0.72 secs – a speed-up of 20x.
- ❑ The results of both versions are numerically identical!

Floating point numbers & IEEE 754

Another gradual underflow example:

- ❑ Cholesky factorization of a sparse matrix
- ❑ runtime: 90+ secs (39 secs user, 50 secs system)
- ❑ this example suffered from gradual underflows
- ❑ no possibility to recompile
- ❑ solution: add a small number ($1e-12$) to all matrix elements
- ❑ new runtime: < 9 secs – no system time overhead!

Floating point numbers & IEEE 754

- ❑ Events that can cause (hardware) traps:
 - ❑ division by zero
 - ❑ working with NaNs (Not A Number) – but some applications rely on that, e.g. for missing data points
 - ❑ gradual underflow
- ❑ Those traps can be a performance killer!
- ❑ BTW: Adobe Flash's floating data type initializes the value to NaN!

Large pages

- ❑ The TLB (or address cache) revisited:
 - ❑ mapping capacity: # TLB entries * page size
 - ❑ example (US-IV): 512 entries of 8 kB => 4MB
- ❑ Use large pages (if supported) if your application suffers from many TLB misses:
 - ❑ recompile & link with `-xpagesize` (Sun Studio)
 - ❑ Solaris:
 - ❑ `ppgsz` command (man `ppgsz`)
 - ❑ MPSS preload library (man `mpss.so.1`)


Large pages

- ❑ Changing the pagesize:
 - ❑ use one of the mechanism above
 - ❑ e.g. set pagesize to 4MB
 - ❑ mapping capacity: 512 entries of 4MB => 2GB!
- ❑ Check: with the 'pmap -sx <PID>' command
 - ❑ look for 'heap', 'anon' and 'stack' in the output
 - ❑ example: `pmap -sx 12345 | egrep 'heap|stack|anon'`
 - ❑ Note: although large pages are supported, they might not be available at run time!

Large pages

Effect of large pages: Matrix times vector

❑ column version (bad version)



Two runs:

a) matvec

b) ppgsz -o stack=4M,heap=4M matvec
same binary – no recompilation!

US-VI+ @ 1800 MHz

L1 : 64 kB

L2 : 2 MB

L3 : 32 MB

Peak : 3.6 Gflop/s

Large pages in Linux

- ❑ In Linux 'large pages' are called 'huge pages'
- ❑ This seems still under development
- ❑ Check the availability: `'cat /proc/meminfo'`
- ❑ look for `hugetlbpage.txt` in the kernel sources

Summary

- ❑ You have now heard about
 - ❑ tuning techniques
 - ❑ tools: compilers, analysis tools
 - ❑ libraries
 - ❑ other performance parameters
 - ❑ debuggers: try Totalview
- ❑ Now you have to apply that and get experience!
- ❑ But never forget:

Correct code has the highest
priority – not speed!