

Principles of Programming
for Environmental Engineers
The GNU/Linux Command-line and
Programming in Bash
© Hugo Connery, 2014
Technical University of Denmark

LICENSE

This work is published under the Creative Commons Attribution-ShareAlike 3.0 Unported License.
See: <http://creativecommons.org/licenses/by-sa/3.0/legalcode>

Table of Contents

1	About this Document.....	4
2	Introduction.....	5
2.1	Community.....	5
2.2	History: UNIX, GNU/Linux and Hackers.....	5
2.3	Advice / Warning	5
2.4	The End Result	5
3	Putty and Access.....	6
3.1	Download and Install.....	6
3.2	Installation and Configuration.....	6
3.3	Command-Line	6
3.4	Security and Stability	7
4	Help and Basic Commands.....	7
4.1	Manual Pages	7
4.1.1	The Crossroads.....	7
4.2	Essential Commands	8
4.2.1	Listing the contents of directories:.....	8
4.2.2	Changing directories:.....	8
4.2.3	Using 'Tab' for path completion	9
4.2.4	Making directories.....	9
4.2.5	Removing directories.....	9
4.2.6	Renaming files.....	9
4.2.7	Copying files.....	9
4.2.8	Deleting (Removing) files/directories.....	10
4.2.9	Changing the 'mode' of a file.....	10
4.3	Shell Parsing: the problem of spaces (and other special characters).....	10
4.3.1	Escaping.....	11
4.3.2	Quoting.....	11
4.3.3	Exercise for the Curious.....	11
5	Text File Editing.....	12
5.1	Nano.....	12
5.2	Vi / Vim.....	12
5.3	Emacs.....	13
6	Command-line Editing	13
6.1	Simple version.....	13
6.2	The Vi version.....	13
7	Cut / Paste: The Putty mechanism.....	13
8	The Environment	14
8.1	“Special” variables.....	14
9	Streams	15
9.1	Basic Stream Usage	15
9.2	Pipelines	16
9.2.1	Basic Example.....	16
9.2.2	Convert Input to Arguments: xargs.....	17
9.2.3	Common and Practical Examples.....	17
10	Processes and Signals.....	18

Principles of Programming: Bash

10.1	Signals.....	18
10.2	Listing Processes.....	18
10.2.1	ps.....	18
10.2.2	Jobs.....	19
11	Processes	20
11.1	A Simple Example in the Shell	20
11.2	Changing the Environment.....	20
11.2.1	Single Command Environment Changes.....	20
11.2.2	Permanent Environment Changes.....	21
11.3	Permanently Changing your Shell itself.....	21
11.4	Parallelisation and Process Control.....	21
11.5	Hang Up: Protecting you Long Time Processes	22
11.6	Deliberately throwing away information	22
11.7	Being Nice: running a process deliberately slowly.....	22
12	Programs, Scripts and File Types.....	24
12.1	Determining the “Type” of a File.....	24
12.2	Compiled Programs vs. Scripts.....	24
12.2.1	A little background and history.....	24
12.2.2	Interpreters.....	25
12.3	Making Scripts.....	25
12.3.1	Shell Scripts: BASH.....	25
12.3.2	Python scripts.....	26
12.3.3	Perl, R, Ruby,	26
13	Useful Commands Index.....	27
14	Programming in Bash.....	28
14.1	About.....	28
14.2	Environment and Variables.....	28
14.2.1	Export and sub-shells.....	28
14.2.2	Command-line Arguments.....	29
14.3	Conditional Branching.....	29
14.3.1	if then else.....	29
14.3.2	case.....	30
14.4	Looping.....	30
14.4.1	for.....	30
14.4.2	while.....	30
14.5	Functions.....	31
14.5.1	Libraries.....	31
14.6	Template.....	31
15	Appendix A: Communities.....	32
15.1	DTU Environment.....	32

1 About this Document

This document is a course material for introducing people to using *NIX, and particularly the GNU/Linux command-line (i.e the bash shell).

It was developed to further research in Environmental Engineering at the like named department of the Technical University of Denmark.

The document talks about using Linux on the command-line, and makes no mention of environmental engineering/modeling practice. This is the potential work of the target audience, rather than the subject under discussion.

2 Introduction

This course is about learning how to use a Linux based system for computational work. It is a crash course, which means that:

- you *will* have more questions than answers during the course, which is expected
- you will often be confused *during* the course, but it is hoped that you will develop methods to resolve your confusion. The primary methods proposed are Read The Manual (RTM) and consult the community (in that order).

2.1 Community

The most effective method of self-help is to find a community with which one can interact to learn and contribute. See Appendix A.

2.2 History: UNIX, GNU/Linux and Hackers

GNU/Linux is a derivative of UNIX (of which there are many), which are derivatives of MULTICS. There are two main branches of UNIX, BSD (Berkeley Software Distribution) and System V. All GNU/Linux and UNIX type systems are written *NIX in this document.

*NIX were written by Hackers. This term is eternally mis-used by the media. Hackers are those who exhibit playfully cleverness. The people who break into other people's computers and cause chaos are called Crackers. For a pleasant and eminently readable history of *NIX and the hackers, see:

<http://www.catb.org/~esr/writings/homesteading/hacker-history/>

Hereafter GNU/Linux is written merely as Linux. In truth, Linux is a kernel, with which you can do almost nothing. It is the GNU part that makes it useful. (Topic for research).

2.3 Advice / Warning

If you're wishing to use Linux then you are almost certainly in for a new experience in computing. If all you have ever used are graphical user interfaces (GUI's), then that new experience is going to leave your 'mouse hand' lost and wondering what to do. The keyboard is your assistant, your mind is the master, and the mouse is hiding in a hole (cat present).

2.4 The End Result

*NIX is a steep learning curve for a person who is only used to GUI's.

After this course, and a few months of practice and frustration, but with the help of the community, you will be able to achieve results on Linux using larger datasets than you can on Windows. You will learn that you need to tell Linux *exactly* what you want it to do, because it will do exactly what you ask it to do with all of your privileges.

You will have precise control over what you want done.

3 Putty and Access

To access the Linux modeling systems you need to use the ssh protocol. A nice tool for doing that from Windows is called Putty.

3.1 Download and Install

Download from:

<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>

Download: A Windows installer for everything except PuTTYtel

Run the installer: Next, ..., Install

(Its an open source tool. It will not install nasty toolbars on your browser or do anything else nasty; Next, Next ...).

3.2 Installation and Configuration

Start it (Start menu): Putty, Putty

1. Add to Host Name '<username>@<the-server>' (e.g abcd@hpc-fe.gbar.dtu.dk)
2. Click on Window and change the Row to 40 (or more than 24)
3. Expand 'Window' to 'Colors' and select 'Use system colors' (this gives you a white background)
4. Back to 'Session'
5. In 'Saved Sessions' enter a name for this connection (e.g hpc-fe)
6. Click 'Save'
7. Click 'Open'
8. Enter password
9. ... do stuff on the Linux server ...
10. Type 'exit' to close the connection

Note that Putty will probably be your access mechanism for Linux. Get to know Putty. Customise it to be as comfortable as possible.

3.3 Command-Line

Welcome to your interface to the Linux system. Why don't we use a GUI? Because it takes up lots of memory that we would rather use for the actual work! A GUI would not help you much at all, because the majority of software that you will be using, particularly for parallel processing, is better used on the command-line.

The following sections of this Crash Course will help you to use this command-line.

Your mouse is going to be next to useless here, so get used to loving your keyboard.

(If you never learned to touch type ...)

3.4 Security and Stability

You may feel uncomfortable, but you need not feel afraid. Whatever you do, you cannot kill the system or damage anyone else's things (unless you try really hard).

4 Help and Basic Commands

4.1 Manual Pages

Every general command-line tool that you will use has a help page, telling you what it does and how to use its command-line switches and arguments to tell it to do exactly what you want.

The 'help' pages in *NIX speak are called Manual Pages and can be accessed with the 'man' command. Try:

```
man man
```

That is, show me the manual page for the 'man' command which shows manual pages. Look at the manual page in structure.

Each manual page comes with sections, and the sections are fairly consistent across Manual Pages (generally known just as man pages).

NAME, SYNOPSIS, DESCRIPTION, ...

Man pages are 'terse'. i.e they say what they mean in the least number of words possible. Reading man pages requires patience, but with a little experience becomes much easier, and you may end up being happy not having to read a book rather than 5 pages.

To start understanding the man pages that you will use, do:

```
man 1 intro
```

The '1' in the above command tells man which section of the manual you wish to look at. Section 1 is for all the command-line tools, in which you are mostly interested.

4.1.1 The Crossroads

At this point most people will have looked at `man man` and thought “how the **** is that going to help me?”.

This is the fork in the road:

1. “I cannot be bothered to 'waste' my time reading this complicated stuff and will ask other people to give me the solution to the specific problem that I to need solve. I am under pressure and do not believe that investing my time in understanding this antiquated insanity is to my, or anyone else's, benefit.”
2. “Crickey, not (Organic Chemistry 101, Calculus and Linear Algebra 101, ... <insert your nightmare course here>) again!” Sadly, yes. Know that the people who wrote these pages are mostly the *authors* of the tool; i.e they **are** the expert. They almost always list their contact addresses, and respond to **well formed questions**.

This is a new world, and this is its initiation. Battle on! Like inorganic chemistry or calculus, the investment of time will be intensely rewarded, though it may seem lightyears away at the outset.

Read The Manual (**RTM**) and then ask questions.

Say, command 'foo' is interesting:

- `man foo`
- `foo -h`
- `foo --help`

Nothing there? Go to *your community*. The web is the 'catch all' (last resort).

4.2 Essential Commands

This section introduces some 'commands' that you will certainly need. In all cases, you should **RTM** for the command to understand it better. Please read the man page about a command before asking others about how to use the command. Failure to understand a man page is part of learning – ask for help. Failure to *read* it is disrespectful.

4.2.1 Listing the contents of directories:

```
ls
ls -l (list in (long) detail)
ls -a (show everything; normally any file/dir that starts with the character '.' is not shown)
```

4.2.2 Changing directories:

```
cd
cd <directory>
```

Relative vs. Absolute path names

'.' is where you are

'..' is the directory above

'/' is the root (start of all file paths)

'/' is the directory separator in *NIX ('\ for windows).

Examples

You are the 'abcd' user, and you are in your home directory `/home/abcd`:

Path	Explanation
<code>foo/bar.txt</code>	The <code>bar.txt</code> file in the <code>foo</code> directory under where you are
<code>./foo/bar.txt</code>	The explicit (relative path) version of the above

<code>/home/abcd/foo/bar.txt</code>	The absolute path to the same file
<code>/tmp/baz.pdf</code>	Absolute path to <code>/tmp/baz.pdf</code>
<code>../../tmp/baz.pdf</code>	The relative path to <code>baz.pdf</code> from the home dir that was assumed above.

4.2.3 Using 'Tab' for path completion

While typing in a file path on the command-line you can use the Tab key to ask the shell to 'finish' the path:

- if you type 'Tab' and there is only one possible completion, the completion is done
- if you type 'Tab' and nothing happens, then there are more than one possible completions, so type 'Tab' again (quickly) and the shell will show you all of the possible options. Add more characters by typing and with another 'Tab' effort the shell will do its best from the new position.

4.2.4 Making directories

```
mkdir <path>
```

4.2.5 Removing directories

```
rmdir <path>
```

will complain if the directory is not empty. See 'rm' for a forceful way to remove entire directory trees (careful!).

4.2.6 Renaming files

Renaming in *NIX is known as 'moving'.

```
mv <old-name> <new-name>
```

The names are paths, thus you can do:

```
mv foo/bar.txt /tmp
```

and it will move the file `foo/bar.txt` into the `/tmp` directory, becoming `/tmp/bar.txt`

4.2.7 Copying files

```
cp <original> <new-file>
```

Recursion

In `cp` and `rm` there is a 'recursive' flag, meaning do this for **everything** under each argument. The flag is `-r`. Read the man page!!!

```
cp -r foo /tmp
```

Copy all of the `foo` subdirectory into `/tmp`.

4.2.8 Deleting (Removing) files/directories

```
rm <file> ...
```

```
rm -r <directory> ...
```

Be very careful with your `<path>`

4.2.9 Changing the 'mode' of a file

In *NIX you can change files to protect them from being changed or deleted, or you can even make them executable.

Make read-only

```
chmod a-w <file> ...
```

Make executable

```
chmod a+x <file> ...
```

(`man chmod`)

You can, of course, do this to entire directory trees with the 'recursive' switch (for `chmod` it is a capital R, because `-r` already has a meaning for `chmod`):

```
chmod -R a-w <dir> ...
```

Doing these things may not seem important now, but will likely be useful later.

4.3 Shell Parsing: the problem of spaces (and other special characters)

The shell's job is to run processes for you. That generally means an executable file and some arguments to be given to it. Thus, the shell needs to be able to work out what the arguments are as opposed to the command itself.

The standard shells, by default, treat any *space* or *tab* as something that separates **words**. The command-line is chopped up into words. The first word in the command-line is the command (to be executed), and the other words are its *arguments*.

This creates potential problems. Say you have a file 'my holiday.txt' and you want to edit it with an editor 'vi' (see below on Text File Editing). You might type this:

```
vi my holiday.txt
```

The shell will see that and chop it up into **three** words: `vi`, `my`, `holiday.txt`

Thus, it runs 'vi' and gives it two arguments, 'my' and 'holiday.txt'. But you wanted to edit the single file 'my holiday.txt'.

There are two basic ways to solve this problem: escaping and quoting

4.3.1 Escaping

If you place a `\` character before **any** character you are telling the shell to treat the next character as an exact literal: no special meaning is to be applied to it. Thus:

```
vi my\ holiday.txt
```

will edit the file because the `' '` (space character) has had it's normal meaning (i.e it separates words) removed.

If you are using 'Tab' for auto-completion, you will see that during the auto-completion, the shell actually puts the `\` in there for you !!!

4.3.2 Quoting

You can group many characters into one **word** by placing single, or double quotes around the characters. E.g

```
vi 'my holiday.txt'
```

or

```
vi "my holiday.txt"
```

In the above two examples there is no difference because the string `"my holiday.txt"` contains only one type of 'special' character, the space symbol. But there are other special symbols, the most common for the shell you are using are each one of:

```
! $ # % & ( ) * ?
```

For the single quoted string `'my holiday.txt'`, the single quoting means: treat the entire string as just plain text – no special meaning *anywhere*. For the double quoted string `"my holiday.txt"` it tells the shell to look in there for **some other** special characters. This is very useful in advanced use of the shell which is beyond the scope of this course.

In summary, use single quotes for any string that you want to be grouped into a single **word** (argument) and all will be good.

4.3.3 Exercise for the Curious

Do the following, exactly, on a command-line and watch what is printed out for the filename completion at the end:

```
mkdir foo
cd foo
touch 'a$!#&*(){}?.rubbish'
ls a<Tab>
```

The completion should look something like: `a\$\! \#\&\(\)\{\}\? .rubbish`

The shell is telling you that these symbols have a special meaning to it.

Then, clean up:

```
cd ..
rm -rf ./foo
```

5 Text File Editing

5.1 Nano

Nano is a trivial text editor. It works exactly like notepad/wordpad from Microsoft. Arrow keys to move around and backspace/delete or typing. See the bottom of the screen. You use Control characters for saving and quitting.

5.2 Vi / Vim

Vi is a stalwart of the *NIX world. It is installed by default on every *NIX, even under a minimal install.

I could wax lyrical about the wonders of vi for hours. But, the essential understanding is that it is a *modal* editor. To understand that, consider nano/notepad etc.. In these editors whenever you type a letter it appears on the screen. Vi has two modes, input mode (when letters appear on the screen) and command mode, when letters do other things, like move you around, delete things, start an input mode etc.. To move out of input mode and into command mode you press Esc.

Here is the bare minimum of the commands in command mode:

Command	Description
h, j, k, l	Move left, down, up, right (respectively)
W, w, B, b	Move forward: one word (ignoring funny characters), one word (honoring funny characters), and back one word (ignoring and honoring) (respectively)
(,), {, }	Move one sentence back, forward, or one paragraph, back, forward
I, i, A, a, O, o	Insert at the start of line, or here. Append to the end of line, or here. Start a new line above where you are, or below.
d<movement>, dd	Delete text to wherever the move takes you, delete this line.
y<movement>, yy	Copy to wherever the move takes you. Copy this line.
P, p	Paste before where you are, paste after where you are.
U, u	Undo all changes on this line, undo the last change (repeat for more undo)
X, x	Delete the character before the cursor, under the cursor
:e <path>	Edit the file <path>
:e #	Edit the file you were editing last
:n	Edit the next file given on the command line
:w	Save this file
:q	Quit
:wq	Save then quit
:q!	Quit and don't save anything. Just quit.
:w!	Save and don't warn me, just do it.

If you are going to be using *NIX for years, then learn something better than nano. Vi is an excellent choice.

The 'Vim' from the title is 'Vi iMproved'. An excellent piece of work. Vim is now the preferred standard vi.

5.3 Emacs

Emacs is the other great *NIX editor. It is colossal. You can run a web-server with it if you are so inclined. The author has no experience with it.

Historical Note:

There are two camps in the *NIX world when it comes to text editors. The vi camp, and the Emacs camp. They love to tell each other that their editor is better. Its a bit like the Windows camp vs. the Mac camp. Religious war. :-)

6 Command-line Editing

6.1 Simple version

Use the Up and Down arrows to browse through your previous commands. Use Left and Right to move and type or Backspace/Delete.

6.2 The Vi version

```
set -o vi
```

Now your command line editing is in vi mode. Movement and editing as per vi (j, k for up, down etc.). Note you start in 'insert' mode after typing return (on the previous command). Thus, if you want to move around (k,j,w etc.) you start with 'Esc' to get into command mode.

7 Cut / Paste: The Putty mechanism

When you high-light text with the mouse it is *automatically* copied. Click the right mouse button to 'paste' it.

Be careful with this. If you paste on the command-line, the shell will just start doing whatever the paste told it to do.

8 The Environment

In a *NIX system every process has access to a dictionary of name / value pairs. A process can read what is already there and use that information, or change what is there, or even delete what is there or create new name / value pairs.

Many 'programs' use the environment as an alternative, or in addition to the command-line switches that we have seen.

Over time, as some information is equally useful to many programs, these values have been preferentially read from the environment. This is because the environment is trivial for a bunch of programs to share. More on this in Process handling.

Some examples of environment name/value pairs:

USER=abcd (your user name)

HOME=/home/abcd (where your home directory is)

One can see how these values can be useful for many programs.

By convention, all environment variables are written in UPPER CASE.

To see all of the environment variables that are set in your shell do:

```
env
```

8.1 “Special” variables

There are a small collection of very useful environment variables that the shell itself maintains for you. See 'Special Paramters' in the bash man page. Some are listed below. Note that the man page is **the** reference.

Variable	Meaning / Comment
\$?	The exit status of the last command executed. One uses this all the time.
\$@	The arguments to a script, all together, as words.
\$0, \$1, ...	The individual parameters from the command line. 0 is the script itself, 1 is the first argument etc.
\$#	The number of arguments, not including the script itself. E.g if \$# is 3 then \$3 will have a value but \$4 will not.
\$\$	The process id of the current process (i.e the shell).

9 Streams

A stream is a sequence of bytes, followed eventually by an 'end-of-file' marker (meaning that the stream is at its end). A file is a stream. But, a stream can exist just in memory as two (or more) processes communicate with each other, and thus avoids the HEAVY overhead of using a file. A stream could be seen as an ephemeral file (of the other way around; a file is a stream stored on a disk ;-).

In *NIX each and EVERY process (program) gets given 3 streams by default:

Stream	Usage	Number	Explanation
Stdin	<	0	The standard input. That is where a program will read from, by default.
Stdout	>	1	The standard output. That is where a program will write data, by default.
Stderr	2>	2	The standard error. This is a where a program will write informational messages, or error notifications, by default.

9.1 Basic Stream Usage

If you have a command that prints out information, and you want to save it, you just tell the shell to re-direct its stdout to a file:

```
cmd > a-file.txt
```

If a command reads the information on which it will work from the stdin, then you can re-direct its stdin from a file:

```
cmd < b-file.txt
```

If you just want to see what a command prints out, but want to save any *messages* it prints into a file, you can re-direct its stderr:

```
cmd 2> c-file.txt
```

If you use a double greater than (>>) then the information printed will be appended to the file, if it already exists.

Of course, all of the above can be combined:

```
cmd < a-file.txt > b-file.txt 2>> c-file.log
```

Here we have re-directed all standard streams, but with stderr writing in append mode (>>). Note that the names of the files are purely for example. They can be anything you want (remember the 'space character' problem! You need to quote your file names if you wish to use spaces in them.)

Finally, if you want to redirect both stderr and stdout to the same file, there is a way of doing that too. The order in which you write this is important. Here is how to say: redirect stdout to a-file.txt and also redirect stderr to wherever stdout is going:

Principles of Programming: Bash

```
cmd > a-file.txt 2>&1
```

The `&1` means 'wherever 1 is going' (1 is stdout, remember?). Similarly, but written a different way: redirect stderr to a-file.txt and send stdout to wherever stderr is going:

```
cmd 2> a-file.txt >&2
```

If we get the order wrong, things will not be what you expect. Say, we take the first example and write it in the other order:

```
cme 2>&1 > a-file.txt
```

Think carefully about this. It says, redirect stderr to wherever stdout is going. At this point, stdout is going to the terminal (window), so stderr is sent there too (it was already going there, but that's what you said to do). Then stdout is redirected to a-file.txt.

This is a bit classic *NIX. It does **exactly** what you asked it to do, even if that is not what you meant.

All of these 'tricks' are mildly useful. But where it really ROCKS is a pipeline.

9.2 Pipelines

Some command do things and print their results to stdout and other commands read the things on which they should work from stdin.

Wouldn't it be nice to be able to connect the first with the second?

One can do:

```
cmd-A > file.txt
cmd-B < file.txt
rm file.txt
```

All command-line tools use at least some of these streams, as it creates a standard manner in which you can control them, and additionally allows them to easily work together.

The above is lots of typing and we're wasting time with the file system.

A pipeline is connection of at least two commands together separated by the `|` symbol (called 'pipe' in *NIX language, for obvious reasons).

What the pipe symbol does is to tell the shell that you want the stdout of the command before the pipe symbol connected to the stdin of the one after. I.e cmd-B will read what cmd-A printed.

To make that pipeline:

```
cmd-A | cmd-B
```

And that's it. You can make these pipelines as long as you want: `cmd-A | cmd-B | cmd-C | ...`

9.2.1 Basic Example

Say you want to know the number of files in a directory which end in .txt?

You can:

```
ls *.txt
```

and then count them all. But, there is a command that counts characters, words or lines for you called 'wc' (word count). If we use `ls -l *.txt` then ls will print out each file name on a

different line. Then we can use 'wc -l' (count lines). And so, to count these files we just use:

```
ls -l *.txt | wc -l
```

9.2.2 Convert Input to Arguments: xargs

Say we want to know the number of lines in each of the text files in the directory, and we want that sorted by the number of lines? I.e we get a table of numbers in the first column and file names in the second, sorted by the number of lines.

The problem here is that we want the output of the 'ls' command to become arguments for the wc command (read its man page). There is a special tool for that called xargs. Here's the first part of the solution:

```
ls -l *.txt | xargs wc -l
```

Now wc will count the lines in each *file* that ls printed out because xargs is reading what ls printed out and then putting them on the command-line as arguments for wc.

Then, to sort them, we just use the sort command (with -n for sort numerically):

```
ls -l *.txt | xargs wc -l | sort -n
```

Now you have a list of the number of lines in each *.txt file in the directory sorted by the number of lines.

9.3 Common and Practical Examples

Say we want to count the number of files below a directory:

```
find <path to directory> -type f | wc -l
```

Say we want to find the number of files that match a pattern in their name below a directory:

```
find <path> \( -type f -a -name '*.txt' \) | wc -l
```

9.4 Regular Expressions: grep and sed

Every command-line environment from any useful operating system comes with a way to match file named. You may know of *.* from DOS, meaning match anything of any length followed by a dot followed by anything of any length. The file name pattern matching rules are generally known as Globbs. We have those for Bash too (RTM).

*NIX introduced a far more powerful text pattern matching mechanism called Regular Expressions. Regular Expressions are commonly shortened to 'regex' or 'regexp'. With the additional power comes a little complexity. However, when you learn how to use just 20% of that you immediately get two of the most powerful tools: grep and sed.

9.4.1 grep

The best complete introduction to grep (and regular expressions) is maintained by GNU:

<http://www.gnu.org/software/grep/manual/>

The best way to learn this is to play with it on the command-line. You come up with an expression and test it using echo.

```
echo some-text | grep test-expression
```

Principles of Programming: Bash

Read the guide above and here are some examples:

Match just simple characters:

```
prompt$ echo abcde | grep bc
abcde
prompt$ echo abced | grep cb
prompt$
```

Note that `cb` does not occur in `abcde`. Order matters.

```
prompt$ echo abcde | grep B
prompt$
```

So does case!

Using character classes:

```
prompt$ echo abcde 12345 | grep '[:alpha:]'
abcde 12345
prompt$ echo abcde 12345 | grep '[:digit:]'
abcde 12345
```

And now with anchors:

```
prompt$ echo abcde 12345 | grep '[:alpha:]$'
prompt$ echo abcde 12345 | grep '^[:digit:]'
```

In these cases we asked for letters (alphabets) and the end of line, and then numbers (digits) at the beginning, which we do not have.

9.4.2 Sed

Sed (Stream Editor) allows one to use regular expressions to match text and then do stuff with it. Most commonly one will replace. So, sed is the the great 'search and replace' mechanism.

Say, you just want to insert a `+` symbol at the beginning of each line in a file.

```
Sed 's/^/+/g' < file > tmp ; mv tmp file
```

Sed is natively built in to Vi too. So to do the above in Vi (whilst in command mode):

```
:%s/^/+/g
```

The `%` above means 'whole of file'. The rest is exactly the same sed argument from above.

10 Processes and Signals

10.1 Signals

*NIX systems have a way of you sending a message to a process at any time you wish. When you do this the operating system itself acts on your behalf to stop the process doing what its doing and asking it to deal with the 'telegram' (signal).

There are many different signals, and how they are used has become standardised by convention over time. Here we talk about only the most commonly used. Why they are important will become increasingly clear in the Processes section.

The core signals are:

Signal Abbreviation	Meaning
---------------------	---------

Principles of Programming: Bash

INT	Interrupt: tell the process to clean up and stop as soon as it has cleaned up
KILL	Kill: terminate the process. It gets no chance to clean up and is just annihilated
USR1	User Configurable 1: Send a signal whose message is entirely dependent on the program to which you send it (read the man page).
USR2	Same as USR1, but the meaning of the symbol may be different. (RTM)
HUP	Hang Up: This comes from the old modem days. It means that the user has just disconnected from the system, and thus, all of the things that they were doing should be informed. By default, each process will treat it as an INT signal (clean up and stop)

Sending a signal is easy, but you have to know which process you wish to send it to. For that, you need the 'ps' or 'jobs' commands.

The operating system will refuse to send signals to processes that you don't own (I.e you did not start them).

To send an INT signal to the currently running (foreground) process, type Ctrl-c. This means 'copy' in Windows land, but **stop** what you are doing now in *NIX land. (Careful)

10.2 Listing Processes

10.2.1 ps

ps is like 'ls' but for processes rather than files. Read the man page!

```
ps          (tell me about my processes )
```

```
ps -ef      (tell me everything about all processes )
```

Read the output. The number in the left most column is the process id (number). This uniquely identifies a process.

10.2.2 Jobs

Listing jobs that your shell is running in the background:

```
jobs
```

They are denoted by a percent sign (%). Thus [1] (as printed by 'jobs') is referred to as %1 when used on the command line to signal a process.

Send a Signal:

```
kill -(SIGNAL NAME) process_id
```

e.g

Principles of Programming: Bash

```
kill -INT 43231
```

or with a 'job':

```
kill -INT %2 (send an INT signal to job two)
```

11 Processes

*NIX has only one way for a new 'program' (process) to be created. It is known as a 'fork' (based on the OS interface that requests a new processes). A fork of an existing process creates an entirely new process that is an exact copy of the one that requested the fork.

That may seem rather useless. However, the new 'child' process can then request (just as any process can, at any time) that it is overwritten by some program that it can find. This is called an 'exec'.

Thus, the normal course of events for a new process is first a fork of the parent process, then an exec of what the new process is meant to do.

The question is, what is NOT changed by the exec? The brief answer is:

- All of the environment is preserved
- All file handles are preserved (i.e the streams)

11.1 A Simple Example in the Shell

What happens when you type 'ls'?

First, understand that for the shell itself is connected to the command-line window. That is, its stdin reads from there, and both its stdout and stderr write there. Think about that.

So, when you type 'ls' (and press Return) the following occurs:

1. the shell does a fork, and *instructs* the child to load the 'ls' programs (exec). Meanwhile the shell itself is WAITING for the child to finish.
2. (we're now talking about the new 'child' process) the child inherits all of the environment and the file handles, and then performs the exec, as instructed.
3. now the 'ls' program takes over in the 'child' process but still has the same inherited environment and file handles. It starts to do what it is asked to do (read the current directory and print out its contents in default formatting).
4. It does this and uses the stdout that it was given, so the 'printing' that it does appears on the command-line window (I.e it 'inherited' that from the shell)
5. (back to the shell) It detects that the child has finished and then shows you another prompt so you can ask it to do something else.

11.2 Changing the Environment

11.2.1 Single Command Environment Changes

If you run:

```
cmd
```

it will get your current environment, run and finish and then you get your prompt again. Simple and normal. Say, you want to change your environment for that command, and that command only?

Lets say that you want 'cmd' to think that you are not user 'abcd' but user 'wxyz'? You do:

```
USER='wxyz' cmd
```

The shell sees this and knows what you mean. So, after it does the fork, and before it does the exec it changes the environment variable USER to the value that you specified. So, when cmd starts it gets a **modified** version of the environment, but nothing changes in the shell itself.

11.2.2 Permanent Environment Changes

There is a way to ask the shell to read a file and allow that file to change its environment permanently. Instead of just typing 'cmd' you prefix it with a dot:

```
. cmd
```

11.3 Permanently Changing your Shell itself

Of course the shell will let you completely overwrite it. This is rarely a good idea, and almost certain to annoy you unless you know exactly what you are doing. Instead of running just 'cmd' you prefix it with 'exec':

```
exec cmd
```

You will not get another prompt unless the command you chose to exec was a shell (there are many of them, of course). This is essentially how you can use another shell.

11.4 Parallelisation and Process Control

After all that mostly techie and seemingly useless information we come to something directly useful. Running a process in the background.

Recall that when you run a command (e.g ls) the shell will fork/exec it and then WAIT for it to finish. What if the command will take lots of time and you want to do other things in the meantime?

Then you instruct the shell to run it, but do it in the *background* and to give you your prompt back immediately:

```
cmd &
```

But, there is a challenge here too. If you do not instruct the background task to do anything specific with its standard streams then they will continue to be associated with your 'window' and thus what is happening in the background gets mixed up with the other things that you are doing.

So, a smarter move is to redirect the stdout and stderr of 'cmd' to a file or two which you can watch or look at when you are ready:

```
cmd > out-file.txt 2> msg-file.txt &
```

Now you can go about your other business whilst 'cmd' goes about its.

And, you can do this again and again:

```
cmd -i foo1 > out-file1.txt 2> msg-file1.txt &  
cmd -i foo2 > out-file2.txt 2> msg-file2.txt &  
cmd -i foo3 > out-file3.txt 2> msg-file3.txt &  
cmd -i foo4 > out-file4.txt 2> msg-file4.txt &
```

Now you have 4 processes running in the 'background'.

As mentioned in the Signals section, you can use 'jobs' to request the shell to tell you about what is happening in the background.

11.5 Hang Up: Protecting you Long Time Processes

As noted in the Signals section, when you disconnect from the computer, every process that you have created and is still running will get told about your disconnection (receive a HUP signal) and is asked to clean up and quit. This is not what you want when you're running a 3 day simulation. But, as with all other things, there is a standard solution to this well known problem.

You need to stop the HUP signal from being delivered to the process! Enter 'nohup'.

All you need to do is to prefix your command with 'nohup'.

So, considering the above listed 4 background processes, we just get:

```
nohup cmd -i foo1 > out-file1.txt 2> msg-file1.txt &
nohup cmd -i foo2 > out-file2.txt 2> msg-file2.txt &
nohup cmd -i foo3 > out-file3.txt 2> msg-file3.txt &
nohup cmd -i foo4 > out-file4.txt 2> msg-file4.txt &
```

Note that nohup expects (and checks) the you have redirected stdout and stderr. That is because when you disconnect you cannot receive any information sent to the 'window' that is now gone (you are disconnecting). Thus, nohup wants you to redirect your stdout and stderr, and will do this for you automatically if you have not.

This behaviour could be restated as '*NIX will not throw away information unless you explicitly ask it to'.

11.6 Deliberately throwing away information

*NIX has a universal garbage bin called /dev/null. It is a file to which you can write eternally and whatever you write there will be immediately discarded. At first glance this seems silly. But, it has several standard uses.

Say, you want to run a program in the background, under nohup, and you don't care what it prints out. Maybe you know that it will create a file and write its data there. Thus you could say.

```
nohup cmd > /dev/null 2> log-file.txt &
```

What 'cmd' prints out (to stdout) is deliberately discarded (sent to /dev/null). We are possibly interested in the data from stderr and have sent that to 'log-file.txt'.

11.7 Being Nice: running a process deliberately slowly

Each process has a priority. Linux uses this to decide what should run next, when there are more processes than CPUs (which is almost always the case). The lower the priority of a process, the more likely it is to get more time actually running.

One can ask for a process to be run slower than the default scheduling. This is done using the 'nice' command (RTM).

Principles of Programming: Bash

This is a nice thing to do, when you have some task that you want done that may take say an hour, but you dont really need it for at least 3 hours. So, just up the nice value and it will have less impact on other people's use of the system. For example,

```
nice -n 10 run_my_calculation
```


12 Programs, Scripts and File Types

12.1 Determining the “Type” of a File

In Windows, the extension of a file indicates the way that that file should be handled. *NIX has a completely different mechanism that places no restrictions on how you name your files. There is no need for a file extension at all.

As there are no naming restrictions, the name, or extension of a file cannot tell you anything certain. To determine what a file is we need to look at it. There is, of course, a tool that will do this for you. It is the 'file' command.

For example, we can ask what type of file the 'ls' program is:

```
file /bin/ls
/bin/ls: ELF 64-bit LSB executable, x86-64, version 1 (SYSV),
dynamically linked (uses shared libs), for GNU/Linux 2.6.32,
stripped
```

The above example tells us that 'ls' is a compiled program (and lots more technical information).

Another example:

```
file /usr/bin/ipython
/usr/bin/ipython: a /usr/bin/python script text executable
```

Here we see that ipython is a 'script text executable' meaning that it is not compiled and can be read by mortal eyes. It also tells us which language interpreter will be used to 'run' it (/usr/bin/python). I.e this is a script written in the Python language.

12.2 Compiled Programs vs. Scripts

12.2.1 A little background and history

The power at the heart of every computer is a Central Processing Unit, or CPU (possibly many of them in modern computers). Each type of CPU knows how to do a fixed number of things. These things are generally extremely low level operations like 'copy the data from this place in memory to that place over there', or 'add these two numbers'. Each type of operation is identified by some number and potentially takes some arguments (like the locations of the memory from which to copy and then write). In essence, these instructions are a series of binary strings. Not letters, just a big string of Ones and Zeros.

Early on in the use of computers the programmers had to actually write the programs in all of these Ones and Zeros. This is an arcane art.

Following that, computer scientists invented languages which were human readable and were designed to express what you wished a computer to do, but in a readable human form. But, they needed a tool to convert the human readable version into the computer readable version. This type of tool is called a compiler. Its job is to perform this conversion from human readable into computer readable. The collection of files which the compiler reads to produce the final computer

readable binary file are called source files, for they are the source (or beginning) of the final 'program'.

Once this 'compilation' is done, one can then run the program as many times as one wants, and when you run it the computer can instantly read and understand what it is being asked to do.

However, if you want to change the program you must change the source and then do the compilation all over again before you can run the changed version to 'see' the change.

12.2.2 Interpreters

The compiler is a translator between some programming language and computer (CPU) speak. Its end result is not instructing the computer, but building a file that contains these instructions.

There is another strategy and that is to translate the source code into instructions for the computer and give those instructions to the computer immediately. This is called an interpreter.

This has a good side and a bad side. The good side is that you don't have to compile the program, you can just 'run' it directly (I.e the interpreter reads it and directly instructs the computer). The bad side is that the instructions that the computer is given are not necessarily as optimised (fast) as they could be, and you have to run the interpreter every time you want to run the program. Interpreting is a rather heavy task, computationally speaking. So, we get convenience but you suffer in speed.

In recent years, computers have achieved two key advantages that now make the use of interpreters more attractive: computers are faster, and they have more memory.

Interpreters have been around since the beginning of human readable computer languages, but they have only become practically valuable in the last decade or two.

12.3 Making Scripts

If you want to make a script (I.e a program that runs directly from the human readable programming language) you, of course, need to understand the language and write the 'code', but you also need to tell *NIX which interpreter it should use to run it. In Windows this is done with the filename extension. In *NIX you use the **first line** of the file to tell it exactly which interpreter to use.

To do that you use the following format on the first line of the file:

```
#!/<path to the interpreter>
```

Examples:

```
#!/bin/bash
```

```
#!/usr/bin/python
```

To run the file, however, you need to tell the system that it is something that we wish to be *able to be run*. To do this you set the executable flag for the file (see chmod above). For a script called 'my-script' you would do:

```
chmod a+x my-script
```

12.3.1 Shell Scripts: BASH

BASH (often just written bash) is the command-line tool that you are using all the time. It is also a programming language! Any command that you type and run on the command-line can also be put

in a file (script) so that you can do it again and again. Indeed, *NIX systems are full of bash type scripts as it is such a useful language for amalgamating/coordinating various related processes/commands.

Read the manual page for bash, and consult the community.

12.3.2 Python scripts

Just like above, and you need to make sure that you set the first line to tell the system which interpreter to use:

```
#!/usr/bin/python
```

12.3.3 Perl, R, Ruby, ...

There are many interpreted languages. Same deal. Set the first line to refer to the correct interpreter.

13 Useful Commands Index

Here is a list of generally useful commands in *NIX (the exact name may vary a tiny bit across all the *NIX's, but are correct for the RedHat type Linux).

In all cases you should read the manual page if the command looks useful.

Command	Example	Very Brief Description
cat	cat a b > c	Concatenate. Print out the contents of a file, or many files and possibly save them in a single merged file.
cd	cd /tmp	Change directory.
chmod	Chmod a+x foo.sh	Change permissions on files. Example adds execution.
cp	cp -a model backup	Copy a file or directory tree
cut	cut -d ';' -f 1 foo	Extract certain columnar data from a file/stream
dd	dd if=/dev/zero of=foo.dat bs=1024 count=10	Direct disk copy. Copy blocks to/from files. The example creates a file of exactly 10 KB of binary zeros.
df	df -B G	File system usage. How much space is use/available on a/all file systems.
du	du -sh .	Tell me how much space a file or complete directory heirarchy is using.
echo	echo "hello, world."	Print out the arguments given.
file	file foo	Tell me what sort of data is in a file
find	find . -type f	Search a directory tree for matching files/directories.
grep	grep foo file.txt	Search for a pattern and print out matches(*)
head	head -n 15 foo	Print out the first part of a file
less	ls -R less	A pager. Hold the data so I can see it all, rather than scrolling off screen.
ls	ls -a	List information about files and directories
mkdir	mkdir bar	Create a directory
mv	mv foo bar	Rename (move) a file or directory
nice	nice -n 10 my_prog	Run a program with less 'importance'
rm	rm foo	Delete (remove) a file.
rmdir	rmdir bar	Remove (delete) and empty directory
sed	sed 's/ ^t/g'	Stream editor: perform automated editing of text (*)
sort	sort -rn foo	Sort data. Dictionary order, numeric, in reverse etc.
tail	tail foo.log	Print out the last part of a file
tar	tar xzvf foo.tgz	Create / Extract / List (from) tar archives. Warning: NEVER create a tar archive with an absolute path name.

Principles of Programming: Bash

		E.g tar cf /tmp/foo.tar /usr/local is WRONG. Instead, go there and use the relative path: cd /usr ; tar cf foo.tar ./local
tee	tail -f foo tee bar	Create a T fitting in a pipe. i.e both print out what is coming in <i>and</i> save it to a file.
which	which ls	Tell me where the program is that you would run.
who	who	Tell me who else is logged in.
xargs	find . xargs ls -ld	Convert what is read, per line, and make it an argument to a command.
zip / unzip	unzip foo.zip	Create / Extract / List ZIP archives

(*) use a common language for the generic matching of text strings, called a Regular Expression. This is commonly abbreviated as 'regex'. Regex's are another extremely powerful arcane art.

14 Programming in Bash

14.1 About

This document provides a quick introduction to the basics of bash shell programming. Herein 'the shell' means the Bourne Again SHell (bash).

A reader should already be confident using bash as an interactive shell and be comfortable with stream redirection and piping, process control, quoting, exit status,

Note that bash has many features and its manual page is *the authoritative reference*. This document only describes the more commonly used elements of its capabilities as used in simple shell programming.

14.2 Environment and Variables

Bash provides a dictionary of key / value pairs known as environment variables. Variables can be created, modified or removed any time:

```
# Create and/or modify with =
cost_AUD='$1.09'
i=0
verb=push
# Remove with unset
unset cost i
```

The value of a variable is accessed by prefixing the key (name) with '\$'

```
echo $verb
echo ${verb}ing
```

The curly braces are used to delimit the key from following text.

14.2.1 Export and sub-shells

Recall that *NIX uses a process tree. Every running process (except the first) has a parent, and processes can create child processes. The shell is no different and allows you to create sub-shells.

The shell limits what its sub-shells receive from its environment. Only variables that have been marked for 'export' have their values copied to the sub-shell, and these are copies so nothing that the sub-shell does can affect the variables in the parent.

Sub-shells are created by surrounding the command sequence by round brackets:

```
( cd /tmp ; ls -lF )
dir=/var
export dir
( cd $var ; ls -lF )
```

14.2.2 Command-line Arguments

When a shell script is invoked (begins running) it receives the arguments from the command-line which invoked it. \$0 has the value of the path to the script itself. \$1, \$2 ... contain the first, second, ... command-line arguments. The collection of arguments after the path to the script (\$1, \$2, ...) can be collectively accessed as a white-space separated string via \$*.

Processing command-line arguments is further discussed below.

14.3 Command-line substitution

14.4 Conditional Branching

The shell has two basic mechanisms for branching:

- if/then/else
- case

14.4.1 if then else

The if/then/else structure chooses a branch of execution based on executing a command and assessing its exit status. The test(1) command is most commonly used to assess conditions. See the man page for test's capabilities. The shell provides 'syntactic sugar' for the test command so that it looks nice. Both of the following do exactly the same thing (are the same thing), that is check that the contents of the dir variable refers to a directory:

```
if test -d $dir
then
    # do something
fi
if [ -d $dir ]
then
    # do something
fi
```

The branching structure must start with an 'if' and can be extended as needed with repeated 'else if' (elif) conditions, and a optional final catch all case 'else':

```
if [ 0 -eq $num ]
then
    #do something
elif [ 1 -eq $num ]
then
```

```
        # do something
elif [ 2 -eq $num ]
then
        # do something
else
        # do something
fi
```

14.4.2 case

The case branching mechanism matches strings with wildcard support, the same that is used for matching filepaths. In its simplest use it is a one way branching mechanism, like if/then/else. However, it also supports conditional and non-conditional cascading. These more advanced uses are not covered here (see the man page).

Case can be used to perform command-line switch matching:

```
case $arg in
'-h'|'-'?'|'-help') # give help message
;;
*) # glob type wildcard catch all (like else)
;;
esac
```

14.5 Looping

Looping is the repeated execution of a sequence of instructions (commands) – the loop body. Note that a loop body may contain any valid sequence of commands, including other loops (thus called sub-loops).

14.5.1 for

The for loop runs the loop body for each word in a list of words, in the order that they are given in the list, with a placeholder variable being given the value of the current word.

```
for foo in bar baz buz
do
    echo $foo
done
```

This would print out each of the three words bar, baz and buz on separate lines.

There is a second form of 'for' which has a completely different focus and is analogous to the for loop in the group of C-like programming languages. See the bash manual page.

14.5.2 while

The while loop has a conditional expression and a loop body. The conditional expression is evaluated, and if its exit status is zero the body is executed, if not the while loop is terminated. When the loop body completes, the conditional is re-evaluated and the process continues (re-execute the loop or end the loop).

Thus:

```
while [ 1 ]
do
    # something
done
```

will continue to execute forever (as test '1' is true – exits with status 0).

To execute a loop a fixed number of times:

```
i=0
max=512
while [ $i -lt $max ] # -lt tests 'less than'
do
    # something
done
```

A common idiom is reading lines from a file:

```
cat a-file | while read line
do
    # something with $line
done
```

Rather than piping a-file you can redirect stdin. Thus, the following is equivalent but more efficient as the 'cat' program is not run:

```
while read line
do
    # something with $line
done < a-file
```

14.6 Functions

The shell supports functions which may be supplied arguments and **return** a numeric value (just like an exit status). They have access to the environment (global namespace) and may also declare local variables.

For example, if one wishes to deliver informational messages to the software user one should send those to stderr. Rather than performing this redirection of echo each time, one can use a simple function

```
function msg()
{
    echo "$*" >&2
}
```

By default the return value of a function is the return value of the last command that it executed. So, the above would return the value that echo returned. One may explicitly choose a value using the 'return' statement:

```
return 1 # error
```

14.6.1 Libraries

The source command can be used to request that a script runs another script within itself (thus allowing the new script to modify the call script's environment and function definitions).

This enables scripts to load libraries of useful software. This can be useful if a collection of scripts have some common functions, for example error messaging, which can then be shared across the suite.

The '.' is an alias for 'source'.

```
source common-functions.sh
```


14.7 Template

The following is a useful template for small scripts:

```
#!/bin/bash
# Purpose:
# Written by:
# Last edited on:

# The name of the script itself
ME=$(basename $0)
# The path to the script (from wherever the user is)
DIR=$(dirname $0)

# Always provide a usage/help
function usage()
{
    echo "Usage: $ME <-a blah> ..."
    $ME digs holes and fills them in again"
}

# simple messaging (to stderr)
function msg()
{
    echo "$ME: $" >&2
}

# simple 'error' type message
function err()
{
    msg "Error: $"
}

# parse args
case $1 in
    '-h'|-'?'|'-help')
        usage ; exit 0
        ;;
    *)
        ;;
esac

status=0
# do stuff here
#
exit $status
```

15 Appendix A: Communities

15.1 DTU Environment

DTU Environment has a community that supports its modeling based research. The email list is:

community-modeling@env.dtu.dk

The administrators of the list are reachable via the admin mailing list:

community-modeling-admins@env.dtu.dk

Please, just ask the admins to add you to the community, and there you are.

We have few experts and many learners, thus it is essential that the learners help each other as the

Principles of Programming: Bash

experts do not have time to help all. It is hoped for, and expected that, the community will hold regular meetings (lunch-time) to discuss common problems, so that they can better help each other, and so that the experts can help the most learners as possible with the best solutions that they can develop.

Your most important resource is:

<http://wiki.model.env.dtu.dk>

This is available from anywhere inside the network at work, or via the terminal servers (dont know what that means? Talk to IT or a colleague.)